

Real Estate

Project Task: Week 1

df_train Import and Preparation:

1. Import df_train.
2. Figure out the primary key and look for the requirement of indexing.
3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

Exploratory df_train Analysis (EDA):

4. Perform debt analysis. You may take the following steps:
 - a) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map. You may keep the upper limit for the percent of households with a second mortgage to 50 percent
 - b) Use the following bad debt equation: $\text{Bad Debt} = P(\text{Second Mortgage} \cap \text{Home Equity Loan})$
 $\text{Bad Debt} = \text{second_mortgage} + \text{home_equity} - \text{home_equity_second_mortgage}$
 - c) Create pie charts to show overall debt and bad debt
 - d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities
 - e) Create a collated income distribution chart for family income, house hold income, and remaining income

Project Task: Week 2

Exploratory Data Analysis (EDA):

1. Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):
 - a) Use pop and ALand variables to create a new field called population density
 - b) Use male_age_median, female_age_median, male_pop, and female_pop to create a new field called median age
 - c) Visualize the findings using appropriate chart type
2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.
 - a) Analyze the married, separated, and divorced population for these population brackets
 - b) Visualize using appropriate chart type
3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.
4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

Project Task: Week 3

Data Pre-processing:

1. The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables.
2. Each variable is assumed to be dependent upon a linear combination of the common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as "specific variance" because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data. Following are the list of latent variables:
 - Highschool graduation rates

- Median population age
- Second mortgage statistics
- Percent own
- Bad debt expense

Project Task: Week 4

Data Modeling :

1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deplotment_RE.xlsx'. Column hc_mortgage_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location. Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc_mortgage_mean.
 - a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.
 - b) Run another model at State level. There are 52 states in USA.
 - c) Keep below considerations while building a linear regression model. Data Modeling :
 - Variables should have significant impact on predicting Monthly mortgage and owner costs
 - Utilize all predictor variable to start with initial hypothesis
 - R square of 60 percent and above should be achieved
 - Ensure Multi-collinearity does not exist in dependent variables
 - Test if predicted variable is normally distributed

Data Reporting:

2. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:
 - a) Box plot of distribution of average rent by type of place (village, urban, town, etc.).
 - b) Pie charts to show overall debt and bad debt.
 - c) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map.
 - d) Heat map for correlation matrix.
 - e) Pie chart to show the population distribution across different types of places (village, urban, town etc.)

Week 1

df_train Import and Preparation:

1. Import df_train.

```
In [1]: # Import df_train & libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #reading the df_train
df_train = pd.read_csv('train.csv')
df_test = pd.read_csv('test.csv')
print("Shape of train df_train:",df_train.shape)
print("Shape of test df_train:",df_test.shape)
```

```
Shape of train df_train: (27321, 80)
Shape of test df_train: (11709, 80)
```

2. Figure out the primary key and look for the requirement of indexing.

```
In [3]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27321 entries, 0 to 27320
Data columns (total 80 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   UID                                    27321 non-null  int64
1   BLOCKID                               0 non-null      float64
2   SUMLEVEL                              27321 non-null  int64
3   COUNTYID                              27321 non-null  int64
4   STATEID                               27321 non-null  int64
5   state                                 27321 non-null  object
6   state_ab                              27321 non-null  object
7   city                                  27321 non-null  object
8   place                                 27321 non-null  object
9   type                                  27321 non-null  object
10  primary                               27321 non-null  object
11  zip_code                              27321 non-null  int64
12  area_code                             27321 non-null  int64
13  lat                                    27321 non-null  float64
14  lng                                    27321 non-null  float64
15  ALand                                  27321 non-null  float64
16  AWater                                 27321 non-null  int64
17  pop                                    27321 non-null  int64
18  male_pop                              27321 non-null  int64
19  female_pop                            27321 non-null  int64
20  rent_mean                             27007 non-null  float64
21  rent_median                           27007 non-null  float64
22  rent_stdev                            27007 non-null  float64
23  rent_sample_weight                    27007 non-null  float64
24  rent_samples                          27007 non-null  float64
25  rent_gt_10                            27007 non-null  float64
26  rent_gt_15                            27007 non-null  float64
27  rent_gt_20                            27007 non-null  float64
28  rent_gt_25                            27007 non-null  float64
29  rent_gt_30                            27007 non-null  float64
30  rent_gt_35                            27007 non-null  float64
31  rent_gt_40                            27007 non-null  float64
32  rent_gt_50                            27007 non-null  float64
33  universe_samples                      27321 non-null  int64
34  used samples                          27321 non-null  int64
35  hi_mean                               27053 non-null  float64
36  hi_median                             27053 non-null  float64
37  hi_stdev                              27053 non-null  float64
38  hi_sample_weight                      27053 non-null  float64
39  hi_samples                            27053 non-null  float64
40  family_mean                           27023 non-null  float64
41  family_median                         27023 non-null  float64
42  family_stdev                          27023 non-null  float64
43  family_sample_weight                  27023 non-null  float64
44  family_samples                        27023 non-null  float64
45  hc_mortgage_mean                      26748 non-null  float64
46  hc_mortgage_median                    26748 non-null  float64
47  hc_mortgage_stdev                     26748 non-null  float64
48  hc_mortgage_sample_weight              26748 non-null  float64
49  hc_mortgage_samples                    26748 non-null  float64
50  hc_mean                               26721 non-null  float64
51  hc_median                             26721 non-null  float64
52  hc_stdev                              26721 non-null  float64
53  hc_samples                            26721 non-null  float64
54  hc_sample_weight                      26721 non-null  float64
55  home_equity_second_mortgage            26864 non-null  float64
56  second_mortgage                       26864 non-null  float64
57  home_equity                           26864 non-null  float64
58  debt                                  26864 non-null  float64
59  second_mortgage_cdf                    26864 non-null  float64
60  home_equity_cdf                       26864 non-null  float64
61  debt_cdf                              26864 non-null  float64
62  hs_degree                             27131 non-null  float64
63  hs_degree_male                        27121 non-null  float64
64  hs_degree_female                      27098 non-null  float64
65  male_age_mean                         27132 non-null  float64
66  male_age_median                       27132 non-null  float64
67  male_age_stdev                        27132 non-null  float64
68  male_age_sample_weight                 27132 non-null  float64
69  male_age_samples                      27132 non-null  float64
70  female_age_mean                       27115 non-null  float64
71  female_age_median                     27115 non-null  float64
72  female_age_stdev                      27115 non-null  float64
```

```

73 female_age_sample_weight      27115 non-null float64
74 female_age_samples            27115 non-null float64
75 pct_own                       27053 non-null float64
76 married                       27130 non-null float64
77 married_snp                   27130 non-null float64
78 separated                     27130 non-null float64
79 divorced                      27130 non-null float64
dtypes: float64(62), int64(12), object(6)
memory usage: 16.7+ MB

```

```
In [4]: df_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11709 entries, 0 to 11708
Data columns (total 80 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   UID                                   11709 non-null  int64
1   BLOCKID                             0 non-null      float64
2   SUMLEVEL                           11709 non-null  int64
3   COUNTYID                           11709 non-null  int64
4   STATEID                             11709 non-null  int64
5   state                               11709 non-null  object
6   state_ab                            11709 non-null  object
7   city                                11709 non-null  object
8   place                               11709 non-null  object
9   type                                11709 non-null  object
10  primary                             11709 non-null  object
11  zip_code                            11709 non-null  int64
12  area_code                           11709 non-null  int64
13  lat                                  11709 non-null  float64
14  lng                                  11709 non-null  float64
15  ALand                               11709 non-null  int64
16  AWater                              11709 non-null  int64
17  pop                                  11709 non-null  int64
18  male_pop                            11709 non-null  int64
19  female_pop                          11709 non-null  int64
20  rent_mean                           11561 non-null  float64
21  rent_median                         11561 non-null  float64
22  rent_stdev                          11561 non-null  float64
23  rent_sample_weight                 11561 non-null  float64
24  rent_samples                       11561 non-null  float64
25  rent_gt_10                         11560 non-null  float64
26  rent_gt_15                         11560 non-null  float64
27  rent_gt_20                         11560 non-null  float64
28  rent_gt_25                         11560 non-null  float64
29  rent_gt_30                         11560 non-null  float64
30  rent_gt_35                         11560 non-null  float64
31  rent_gt_40                         11560 non-null  float64
32  rent_gt_50                         11560 non-null  float64
33  universe_samples                   11709 non-null  int64
34  used_samples                       11709 non-null  int64
35  hi_mean                            11587 non-null  float64
36  hi_median                          11587 non-null  float64
37  hi_stdev                           11587 non-null  float64
38  hi_sample_weight                   11587 non-null  float64
39  hi_samples                         11587 non-null  float64
40  family_mean                        11573 non-null  float64
41  family_median                      11573 non-null  float64
42  family_stdev                       11573 non-null  float64
43  family_sample_weight               11573 non-null  float64
44  family_samples                     11573 non-null  float64
45  hc_mortgage_mean                   11441 non-null  float64
46  hc_mortgage_median                 11441 non-null  float64
47  hc_mortgage_stdev                  11441 non-null  float64
48  hc_mortgage_sample_weight          11441 non-null  float64
49  hc_mortgage_samples                11441 non-null  float64
50  hc_mean                            11419 non-null  float64
51  hc_median                          11419 non-null  float64
52  hc_stdev                           11419 non-null  float64
53  hc_samples                         11419 non-null  float64
54  hc_sample_weight                   11419 non-null  float64
55  home_equity_second_mortgage         11489 non-null  float64
56  second_mortgage                    11489 non-null  float64
57  home_equity                        11489 non-null  float64
58  debt                               11489 non-null  float64
59  second_mortgage_cdf                 11489 non-null  float64
60  home_equity_cdf                    11489 non-null  float64
61  debt_cdf                           11489 non-null  float64
62  hs_degree                          11624 non-null  float64
63  hs_degree_male                     11620 non-null  float64
64  hs_degree_female                   11604 non-null  float64
65  male_age_mean                      11625 non-null  float64
66  male_age_median                    11625 non-null  float64

```

```

67 male_age_stdev          11625 non-null float64
68 male_age_sample_weight 11625 non-null float64
69 male_age_samples        11625 non-null float64
70 female_age_mean         11613 non-null float64
71 female_age_median       11613 non-null float64
72 female_age_stdev        11613 non-null float64
73 female_age_sample_weight 11613 non-null float64
74 female_age_samples      11613 non-null float64
75 pct_own                 11587 non-null float64
76 married                 11625 non-null float64
77 married_snp             11625 non-null float64
78 separated               11625 non-null float64
79 divorced                11625 non-null float64

```

dtypes: float64(61), int64(13), object(6)
memory usage: 7.1+ MB

```
In [5]: print(df_train.shape)
df_train['UID'].nunique()
```

(27321, 80)

Out[5]: 27161

```
In [6]: print(df_test.shape)
df_test['UID'].nunique()
```

(11709, 80)

Out[6]: 11677

After checking the Feature 'UID' we can conclude that it is the primary key. But there are duplicate records with same UID so we need to remove those duplicate records

```
In [7]: df_train.UID.nunique()
```

Out[7]: 27161

```
In [8]: df_test.UID.nunique()
```

Out[8]: 11677

```
In [9]: grp_train = df_train.groupby('UID')
grp_train.size().sort_values(ascending=False).head(60)
# grp_train.get_group(282028)
grp_train.get_group(230058)
```

Out[9]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	...	female_age_mean	female_ag
--	-----	---------	----------	----------	---------	-------	----------	------	-------	------	-----	-----------------	-----------

777	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315	
1623	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315	
17489	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315	
26046	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315	

4 rows × 80 columns

```
In [10]: grp_test = df_train.groupby('UID')
grp_test.size().sort_values(ascending=False).head(60)
# grp_test.get_group(282028)
grp_test.get_group(230058)
```

Out[10]:	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	...	female_age_mean	female_ag
	777	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315
	1623	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315
	17489	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315
	26046	230058	NaN	140	73	6	California	CA	Oceanside	Camp Pendleton North	City	...	19.99315

4 rows × 80 columns

```

In [11]: #Dropping the duplicate values
print("Before dropping the duplicates the total rows are:", df_train.shape[0], "Shape:", df_train.shape)
df_train.drop_duplicates(keep='first', inplace=True)
print("After dropping the duplicates the total rows are:", df_train.shape[0], "Shape:", df_train.shape)

```

Before dropping the duplicates the total rows are: 27321 Shape: (27321, 80)
After dropping the duplicates the total rows are: 27161 Shape: (27161, 80)

```

In [12]: #Dropping the duplicate values
print("Before dropping the duplicates the total rows are:", df_test.shape[0], "Shape:", df_test.shape)
df_test.drop_duplicates(keep='first', inplace=True)
print("After dropping the duplicates the total rows are:", df_test.shape[0], "Shape:", df_test.shape)

```

Before dropping the duplicates the total rows are: 11709 Shape: (11709, 80)
After dropping the duplicates the total rows are: 11677 Shape: (11677, 80)

Now we can set the UID as the index

```

In [13]: # backup
week1_bck_train = df_train.copy()
week1_bck_test = df_test.copy()

```

```

In [14]: df_train.set_index('UID', inplace=True)
df_test.set_index('UID', inplace=True)

```

```

In [15]: df_train

```

Out[15]:	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	primary	...	female_age_mean	f
UID													
267822	NaN	140	53	36	New York	NY	Hamilton	Hamilton	City	tract	...	44.48629	
246444	NaN	140	141	18	Indiana	IN	South Bend	Roseland	City	tract	...	36.48391	
245683	NaN	140	63	18	Indiana	IN	Danville	Danville	City	tract	...	42.15810	
279653	NaN	140	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	tract	...	47.77526	
247218	NaN	140	161	20	Kansas	KS	Manhattan	Manhattan City	City	tract	...	24.17693	
...	
279212	NaN	140	43	72	Puerto Rico	PR	Coamo	Coamo	Urban	tract	...	42.73154	
277856	NaN	140	91	42	Pennsylvania	PA	Blue Bell	Blue Bell	Borough	tract	...	38.21269	
233000	NaN	140	87	8	Colorado	CO	Weldona	Saddle Ridge	City	tract	...	43.40218	
287425	NaN	140	439	48	Texas	TX	Colleyville	Colleyville City	Town	tract	...	39.25921	
265371	NaN	140	3	32	Nevada	NV	Las Vegas	Paradise	City	tract	...	34.45345	

27161 rows × 79 columns

```

In [16]: df_test

```

Out[16]:

	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type	primary	...	female_age_mean
	UID											
	255504	NaN	140	163	26	Michigan	MI	Detroit	Dearborn Heights City	CDP	tract ...	34.78682
	252676	NaN	140	1	23	Maine	ME	Auburn	Auburn City	City	tract ...	44.23451
	276314	NaN	140	15	42	Pennsylvania	PA	Pine City	Millerton	Borough	tract ...	41.62426
	248614	NaN	140	231	21	Kentucky	KY	Monticello	Monticello City	City	tract ...	44.81200
	286865	NaN	140	355	48	Texas	TX	Corpus Christi	Edroy	Town	tract ...	40.66618

	238088	NaN	140	105	12	Florida	FL	Lakeland	Crystal Springs	City	tract ...	53.51255
	242811	NaN	140	31	17	Illinois	IL	Chicago	Chicago City	Village	tract ...	33.14169
	250127	NaN	140	9	25	Massachusetts	MA	Lawrence	Methuen Town City	City	tract ...	43.53905
	241096	NaN	140	27	19	Iowa	IA	Carroll	Carroll City	City	tract ...	45.63179
	287763	NaN	140	453	48	Texas	TX	Austin	Sunset Valley City	Town	tract ...	35.99955

11677 rows × 79 columns

3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

In [17]:

```
filter0 = (((df_train.isnull().sum()/df_train.shape[0])*100)>0)
df_delcols_train = df_train.loc[:,filter0.values]
print(df_delcols_train.isnull().sum()/df_delcols_train.shape[0]*100)
print(" "*41)
filter1 = (((df_test.isnull().sum()/df_test.shape[0])*100)>0)
df_delcols_test = df_test.loc[:,filter1.values]
print(df_delcols_test.isnull().sum()/df_delcols_test.shape[0]*100)
```

BLOCKID	100.000000
rent_mean	0.890983
rent_median	0.890983
rent_stdev	0.890983
rent_sample_weight	0.890983
rent_samples	0.890983
rent_gt_10	0.890983
rent_gt_15	0.890983
rent_gt_20	0.890983
rent_gt_25	0.890983
rent_gt_30	0.890983
rent_gt_35	0.890983
rent_gt_40	0.890983
rent_gt_50	0.890983
hi_mean	0.762122
hi_median	0.762122
hi_stdev	0.762122
hi_sample_weight	0.762122
hi_samples	0.762122
family_mean	0.846802
family_median	0.846802
family_stdev	0.846802
family_sample_weight	0.846802
family_samples	0.846802
hc_mortgage_mean	1.627333
hc_mortgage_median	1.627333
hc_mortgage_stdev	1.627333
hc_mortgage_sample_weight	1.627333
hc_mortgage_samples	1.627333
hc_mean	1.759876
hc_median	1.759876
hc_stdev	1.759876
hc_samples	1.759876
hc_sample_weight	1.759876
home_equity_second_mortgage	1.325430
second_mortgage	1.325430
home_equity	1.325430
debt	1.325430

second_mortgage_cdf	1.325430
home_equity_cdf	1.325430
debt_cdf	1.325430
hs_degree	0.533854
hs_degree_male	0.566989
hs_degree_female	0.629579
male_age_mean	0.544899
male_age_median	0.544899
male_age_stdev	0.544899
male_age_sample_weight	0.544899
male_age_samples	0.544899
female_age_mean	0.592762
female_age_median	0.592762
female_age_stdev	0.592762
female_age_sample_weight	0.592762
female_age_samples	0.592762
pct_own	0.762122
married	0.552262
married_snp	0.552262
separated	0.552262
divorced	0.552262

dtype: float64

BLOCKID	100.000000
rent_mean	1.147555
rent_median	1.147555
rent_stdev	1.147555
rent_sample_weight	1.147555
rent_samples	1.147555
rent_gt_10	1.156119
rent_gt_15	1.156119
rent_gt_20	1.156119
rent_gt_25	1.156119
rent_gt_30	1.156119
rent_gt_35	1.156119
rent_gt_40	1.156119
rent_gt_50	1.156119
hi_mean	0.959150
hi_median	0.959150
hi_stdev	0.959150
hi_sample_weight	0.959150
hi_samples	0.959150
family_mean	1.070480
family_median	1.070480
family_stdev	1.070480
family_sample_weight	1.070480
family_samples	1.070480
hc_mortgage_mean	2.098142
hc_mortgage_median	2.098142
hc_mortgage_stdev	2.098142
hc_mortgage_sample_weight	2.098142
hc_mortgage_samples	2.098142
hc_mean	2.286546
hc_median	2.286546
hc_stdev	2.286546
hc_samples	2.286546
hc_sample_weight	2.286546
home_equity_second_mortgage	1.747024
second_mortgage	1.747024
home_equity	1.747024
debt	1.747024
second_mortgage_cdf	1.747024
home_equity_cdf	1.747024
debt_cdf	1.747024
hs_degree	0.667980
hs_degree_male	0.702235
hs_degree_female	0.822129
male_age_mean	0.659416
male_age_median	0.659416
male_age_stdev	0.659416
male_age_sample_weight	0.659416
male_age_samples	0.659416
female_age_mean	0.745054
female_age_median	0.745054
female_age_stdev	0.745054
female_age_sample_weight	0.745054
female_age_samples	0.745054
pct_own	0.959150
married	0.659416
married_snp	0.659416
separated	0.659416
divorced	0.659416

dtype: float64

Dropping Following features:

-'BLOCKID' as it is having all null values and 'SUMLEVEL' as well as 'primary' feature as it is having only 1 entry.

```
In [18]: df_train.SUMLEVEL.nunique()
```

```
Out[18]: 1
```

```
In [19]: print("Before dropping the feature from train data:",df_train.shape)
df_train.drop(columns = ['BLOCKID','primary', 'SUMLEVEL'], inplace = True)
print("After dropping the feature from train data:",df_train.shape)

print("Before dropping the feature from test data:",df_test.shape)
df_test.drop(columns = ['BLOCKID','primary', 'SUMLEVEL'], inplace = True)
print("After dropping the feature from test data:",df_test.shape)
```

Before dropping the feature from train data: (27161, 79)

After dropping the feature from train data: (27161, 76)

Before dropping the feature from test data: (11677, 79)

After dropping the feature from test data: (11677, 76)

```
In [20]: filter0 = (((df_train.isnull().sum()/df_train.shape[0])*100)>0)
df_delcols_train = df_train.loc[:,filter0.values]
print(df_delcols_train.isnull().sum()/df_delcols_train.shape[0]*100)
print("*****39)
filter1 = (((df_test.isnull().sum()/df_test.shape[0])*100)>0)
df_delcols_test = df_test.loc[:,filter1.values]
print(df_delcols_test.isnull().sum()/df_delcols_test.shape[0]*100)
```

rent_mean	0.890983
rent_median	0.890983
rent_stdev	0.890983
rent_sample_weight	0.890983
rent_samples	0.890983
rent_gt_10	0.890983
rent_gt_15	0.890983
rent_gt_20	0.890983
rent_gt_25	0.890983
rent_gt_30	0.890983
rent_gt_35	0.890983
rent_gt_40	0.890983
rent_gt_50	0.890983
hi_mean	0.762122
hi_median	0.762122
hi_stdev	0.762122
hi_sample_weight	0.762122
hi_samples	0.762122
family_mean	0.846802
family_median	0.846802
family_stdev	0.846802
family_sample_weight	0.846802
family_samples	0.846802
hc_mortgage_mean	1.627333
hc_mortgage_median	1.627333
hc_mortgage_stdev	1.627333
hc_mortgage_sample_weight	1.627333
hc_mortgage_samples	1.627333
hc_mean	1.759876
hc_median	1.759876
hc_stdev	1.759876
hc_samples	1.759876
hc_sample_weight	1.759876
home_equity_second_mortgage	1.325430
second_mortgage	1.325430
home_equity	1.325430
debt	1.325430
second_mortgage_cdf	1.325430
home_equity_cdf	1.325430
debt_cdf	1.325430
hs_degree	0.533854
hs_degree_male	0.566989
hs_degree_female	0.629579
male_age_mean	0.544899
male_age_median	0.544899
male_age_stdev	0.544899
male_age_sample_weight	0.544899
male_age_samples	0.544899
female_age_mean	0.592762
female_age_median	0.592762
female_age_stdev	0.592762
female_age_sample_weight	0.592762
female_age_samples	0.592762
pct_own	0.762122

```

married                                0.552262
married_snp                            0.552262
separated                              0.552262
divorced                               0.552262
dtype: float64
*****
rent_mean                              1.147555
rent_median                            1.147555
rent_stdev                             1.147555
rent_sample_weight                     1.147555
rent_samples                           1.147555
rent_gt_10                             1.156119
rent_gt_15                             1.156119
rent_gt_20                             1.156119
rent_gt_25                             1.156119
rent_gt_30                             1.156119
rent_gt_35                             1.156119
rent_gt_40                             1.156119
rent_gt_50                             1.156119
hi_mean                                0.959150
hi_median                              0.959150
hi_stdev                               0.959150
hi_sample_weight                       0.959150
hi_samples                             0.959150
family_mean                            1.070480
family_median                          1.070480
family_stdev                           1.070480
family_sample_weight                   1.070480
family_samples                         1.070480
hc_mortgage_mean                       2.098142
hc_mortgage_median                     2.098142
hc_mortgage_stdev                      2.098142
hc_mortgage_sample_weight              2.098142
hc_mortgage_samples                    2.098142
hc_mean                                2.286546
hc_median                              2.286546
hc_stdev                               2.286546
hc_samples                             2.286546
hc_sample_weight                       2.286546
home_equity_second_mortgage            1.747024
second_mortgage                        1.747024
home_equity                            1.747024
debt                                   1.747024
second_mortgage_cdf                    1.747024
home_equity_cdf                        1.747024
debt_cdf                               1.747024
hs_degree                              0.667980
hs_degree_male                         0.702235
hs_degree_female                       0.822129
male_age_mean                          0.659416
male_age_median                        0.659416
male_age_stdev                         0.659416
male_age_sample_weight                 0.659416
male_age_samples                       0.659416
female_age_mean                        0.745054
female_age_median                      0.745054
female_age_stdev                       0.745054
female_age_sample_weight               0.745054
female_age_samples                     0.745054
pct_own                                0.959150
married                                0.659416
married_snp                            0.659416
separated                              0.659416
divorced                               0.659416
dtype: float64

```

```

In [21]: df_null_train = df_train[df_train.isna().any(axis=1)]
print("Shape of null df_train:", df_null_train.shape)
df_null_train.head()
df_null_test = df_test[df_test.isna().any(axis=1)]
print("Shape of null df_test:", df_null_test.shape)
df_null_test.head()

```

```

Shape of null df_train: (576, 76)
Shape of null df_test: (322, 76)

```

Out[21]:	COUNTYID	STATEID	state	state_ab	city	place	type	zip_code	area_code	lat	...	female_age_mean	f
UID													
265339	3	32	Nevada	NV	Las Vegas	Winchester	City	89119	702	36.111448	...	33.57247	
287596	451	48	Texas	TX	San Angelo	San Angelo City	Town	76903	325	31.431831	...	21.40298	
250903	25	25	Massachusetts	MA	Cambridge	Cambridge City	City	2139	617	42.359478	...	22.53871	
287557	441	48	Texas	TX	Abilene	Tye City	Town	79607	325	32.423876	...	22.72458	
247510	209	20	Kansas	KS	Kansas City	Kansas City City	City	66104	913	39.171767	...	NaN	

```
In [22]: percent_null_train = (df_null_train.shape[0]/df_train.shape[0])*100
print("% of null in train:",percent_null_train)
percent_null_test = (df_null_test.shape[0]/df_test.shape[0])*100
print("% of null in test:",percent_null_test)
```

Now we know that the percentage of NULL data present in the dataset is very low that is 2.12%, which is very low compared to the dataset.

```
In [23]: df_train.dropna(inplace=True)
df_test.dropna(inplace=True)
```

4. Perform debt analysis. You may take the following steps:

Using SQL

```
In [27]: top 2500
```

2500 rows × 5 columns

Out[28]: 0.99249

```
In [29]: top_2500.second_mortgage.max()
```

Out[29]: 0.43363

In the above steps, we first write include the sqldf function from the pandasql library and then we write the query where we specify that the percent ownership should be greater than 10% and the upper limit of second_mortgage to 50% & sorting the data in descending order of the 'second_mortgage' & limiting the records to 2500.


```
In [30]: import plotly.express as px
import plotly.graph_objects as go
```

```
In [31]: #Using graph objects
fig = go.Figure(data=go.Scattergeo(
    lat = top_2500['lat'],
    lon = top_2500['lng']),)
fig.update_traces(marker=dict(size=5))
fig.update_layout(geo=dict(
    scope = 'north america',
    showland = True,
    landcolor = "rgb(212, 212, 212)",
    subunitcolor = "rgb(255, 255, 255)",
    countrycolor = "rgb(255, 255, 255)",
    showlakes = True,
    lakecolor = "rgb(255, 255, 255)",
    showsubunits = True,
    showcountries = True,
    resolution = 50,
    projection = dict(
        type = 'conic conformal',
        rotation_lon = -100
    ),
    lonaxis = dict(
        showgrid = True,
        gridwidth = 0.5,
        range= [ -140.0, -55.0 ],
        dtick = 5
    ),
    lataxis = dict (
        showgrid = True,
        gridwidth = 0.5,
        range= [ 20.0, 60.0 ],
        dtick = 5
    )
), title='Top 2,500 locations with second mortgage is the highest and percent ownership is above 10 percent
fig.update_geos(
    projection_type="orthographic",
    resolution=50,
    showcoastlines=True, coastlinecolor="RebeccaPurple",
    showland=True, landcolor="LightGreen",
    showocean=True, oceancolor="LightBlue",
    showlakes=True, lakecolor="Blue",
    showrivers=True, rivercolor="Blue")
fig.show()
```

Top 2,500 locations with second mortgage is the highest and percent ownersh

```
In [32]: #using Express
geo_df = pd.DataFrame(dict(lat=top_2500['lat'], lon=top_2500['lng'], subreg=top_2500['second_mortgage']))
fig = px.scatter_geo(geo_df, lat="lat", lon="lon", color="subreg", height= 600, projection = "orthographic",
                    title="Top 2,500 locations with second mortgage is the highest and percent ownership is ab
fig.update_geos(
    projection_type="orthographic",
    resolution=110,
    showcoastlines=True, coastlinecolor="RebeccaPurple",
    showland=True, landcolor="LightGreen",
    showocean=True, oceancolor="LightBlue",
    showlakes=True, lakecolor="Blue",
    showrivers=True, rivercolor="Blue"
)
fig.show()
```

Top 2,500 locations with second mortgage is the highest and percent ownersh



b) Use the following bad debt equation: $\text{Bad Debt} = P(\text{Second Mortgage} \cap \text{Home Equity Loan})$
 $\text{Bad Debt} = \text{second_mortgage} + \text{home_equity} - \text{home_equity_second_mortgage}$

```
In [33]: df_train['BAD_DEBT'] = df_train['second_mortgage'] + df_train['home_equity'] - df_train['home_equity_second_mortgage']
```

Out[33]:

	COUNTYID	STATEID	state	state_ab	city	place	type	zip_code	area_code	lat	...	female_age_median	
	UID												
	267822	53	36	New York	NY	Hamilton	Hamilton	City	13346	315	42.840812	...	45.33333
	246444	141	18	Indiana	IN	South Bend	Roseland	City	46616	574	41.701441	...	37.58333
	245683	63	18	Indiana	IN	Danville	Danville	City	46122	317	39.792202	...	42.83333
	279653	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	927	787	18.396103	...	50.58333
	247218	161	20	Kansas	KS	Manhattan	Manhattan City	City	66502	785	39.195573	...	21.58333

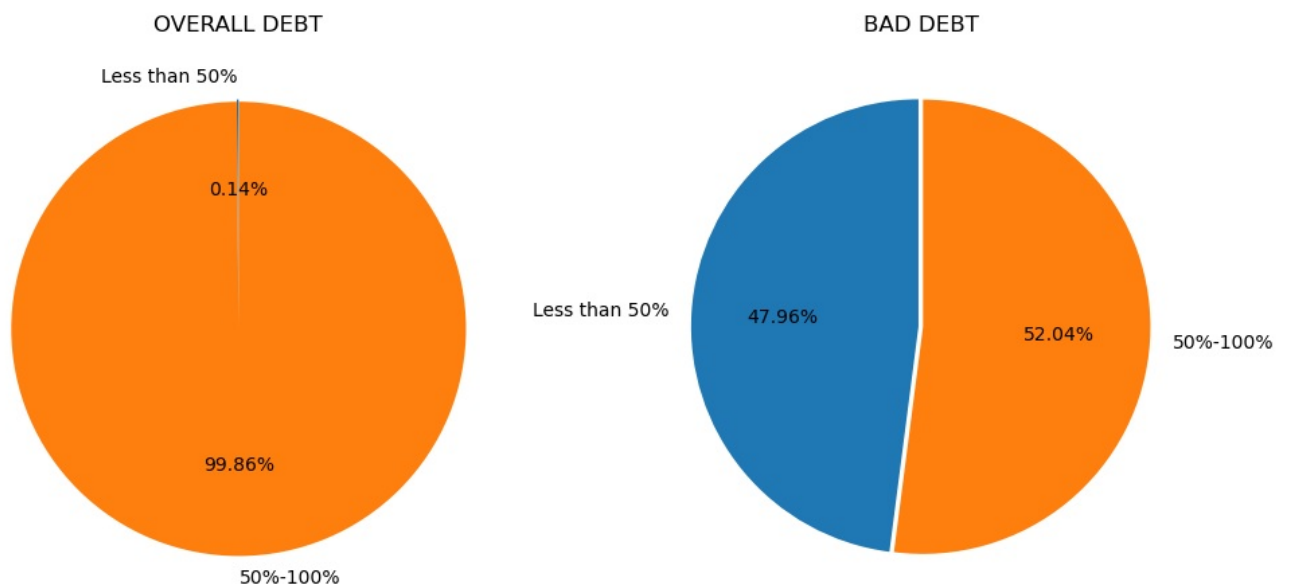
	279212	43	72	Puerto Rico	PR	Coamo	Coamo	Urban	769	787	18.076060	...	40.16667
	277856	91	42	Pennsylvania	PA	Blue Bell	Blue Bell	Borough	19422	215	40.158138	...	39.50000
	233000	87	8	Colorado	CO	Weldona	Saddle Ridge	City	80653	970	40.410316	...	46.33333
	287425	439	48	Texas	TX	Colleyville	Colleyville City	Town	76034	817	32.904866	...	43.41667
	265371	3	32	Nevada	NV	Las Vegas	Paradise	City	89123	702	36.064754	...	29.83333

26585 rows × 77 columns

c) Create pie charts to show overall debt and bad debt

```
In [34]: df_train['BINS_OD'] = pd.cut(df_train['debt'], bins=[0,0.10,1], labels=['Less than 50%', '50%-100%'])
df_train['BINS_BD'] = pd.cut(df_train['BAD_DEBT'], bins=[0,0.10,1], labels=['Less than 50%', '50%-100%'])

pop1=(0,0.01)
pop2=(0.01,0.01)
plt.figure(figsize=(12,10))
plt.subplots_adjust(wspace=0.2)
plt.subplot(1,2,1)
plt.title('OVERALL DEBT')
df_train.groupby(['BINS_OD']).size().plot(kind = 'pie', subplots=True, startangle = 90, autopct = '%1.2f%%', expand=True)
plt.ylabel("")
plt.subplot(1,2,2)
plt.title('BAD DEBT')
df_train.groupby(['BINS_BD']).size().plot(kind = 'pie', subplots=True, startangle = 90, autopct = '%1.2f%%', expand=True)
plt.ylabel("")
plt.show()
```



d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities

```
In [35]: city = df_train.city.value_counts()
```

```
city_top_10 = city.head(10)
city_top_10.index
```

```
Out[35]: Index(['Chicago', 'Brooklyn', 'Los Angeles', 'Houston', 'Philadelphia',
              'San Antonio', 'Baltimore', 'Las Vegas', 'Phoenix', 'Miami'],
              dtype='object')
```

```
In [36]: df_train['GOOD_DEBT'] = 1 - df_train.BAD_DEBT
df_train
```

```
Out[36]:
```

	COUNTYID	STATEID	state	state_ab	city	place	type	zip_code	area_code	lat	...	female_age_samples	
	UID												
	267822	53	36	New York	NY	Hamilton	Hamilton	City	13346	315	42.840812	...	2618.0
	246444	141	18	Indiana	IN	South Bend	Roseland	City	46616	574	41.701441	...	1284.0
	245683	63	18	Indiana	IN	Danville	Danville	City	46122	317	39.792202	...	3238.0
	279653	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	927	787	18.396103	...	1559.0
	247218	161	20	Kansas	KS	Manhattan	Manhattan City	City	66502	785	39.195573	...	3051.0

	279212	43	72	Puerto Rico	PR	Coamo	Coamo	Urban	769	787	18.076060	...	938.0
	277856	91	42	Pennsylvania	PA	Blue Bell	Blue Bell	Borough	19422	215	40.158138	...	2039.0
	233000	87	8	Colorado	CO	Weldona	Saddle Ridge	City	80653	970	40.410316	...	1364.0
	287425	439	48	Texas	TX	Colleyville	Colleyville City	Town	76034	817	32.904866	...	5815.0
	265371	3	32	Nevada	NV	Las Vegas	Paradise	City	89123	702	36.064754	...	1911.0

26585 rows × 80 columns

```
In [37]: new_df = df_train[['city', 'second_mortgage', 'home_equity', 'GOOD_DEBT', 'BAD_DEBT']]
new_df
```

```
Out[37]:
```

	city	second_mortgage	home_equity	GOOD_DEBT	BAD_DEBT	
	UID					
	267822	Hamilton	0.02077	0.08919	0.90592	0.09408
	246444	South Bend	0.02222	0.04274	0.95726	0.04274
	245683	Danville	0.00000	0.09512	0.90488	0.09512
	279653	San Juan	0.01086	0.01086	0.98914	0.01086
	247218	Manhattan	0.05426	0.05426	0.94574	0.05426

	279212	Coamo	0.00000	0.00000	1.00000	0.00000
	277856	Blue Bell	0.02112	0.19641	0.79092	0.20908
	233000	Weldona	0.02024	0.07857	0.92143	0.07857
	287425	Colleyville	0.07550	0.12556	0.85695	0.14305
	265371	Las Vegas	0.01412	0.18362	0.81638	0.18362

26585 rows × 5 columns

```
In [38]: new_df = new_df[new_df['city'].isin(city_top_10.index)]
new_df.city.value_counts()
```

```
Out[38]: Chicago      286
Brooklyn    261
Los Angeles  219
Houston     213
Philadelphia 160
San Antonio  138
Baltimore   128
Las Vegas   123
Phoenix     114
Miami       105
Name: city, dtype: int64
```

```
In [39]: plt.figure(figsize=(12,20))
```

```
plt.subplots_adjust(hspace=0.35)





plt.subplot(4,1,1)
plt.title('CITY VS SECOND MORTGAGE')
sns.boxplot(x = 'city', y = 'second_mortgage', data = new_df)
plt.xlabel("City")
plt.ylabel("Second Mortgage")

plt.subplot(4,1,2)
plt.title('CITY VS HOME EQUITY')
sns.boxplot(x = 'city', y = 'home_equity', data = new_df)
plt.xlabel("City")
plt.ylabel("Home Equity")

plt.subplot(4,1,3)
plt.title('CITY VS GOOD DEBT')
sns.boxplot(x = 'city', y = 'GOOD_DEBT', data = new_df)
plt.xlabel("City")
plt.ylabel("Good Debt")

plt.subplot(4,1,4)
plt.title('CITY VS BAD DEBT')
sns.boxplot(x = 'city', y = 'BAD_DEBT', data = new_df)
plt.xlabel("City")
plt.ylabel("Bad Debt")

plt.show()
```

e) Create a collated income distribution chart for family income, house hold income, and remaining income

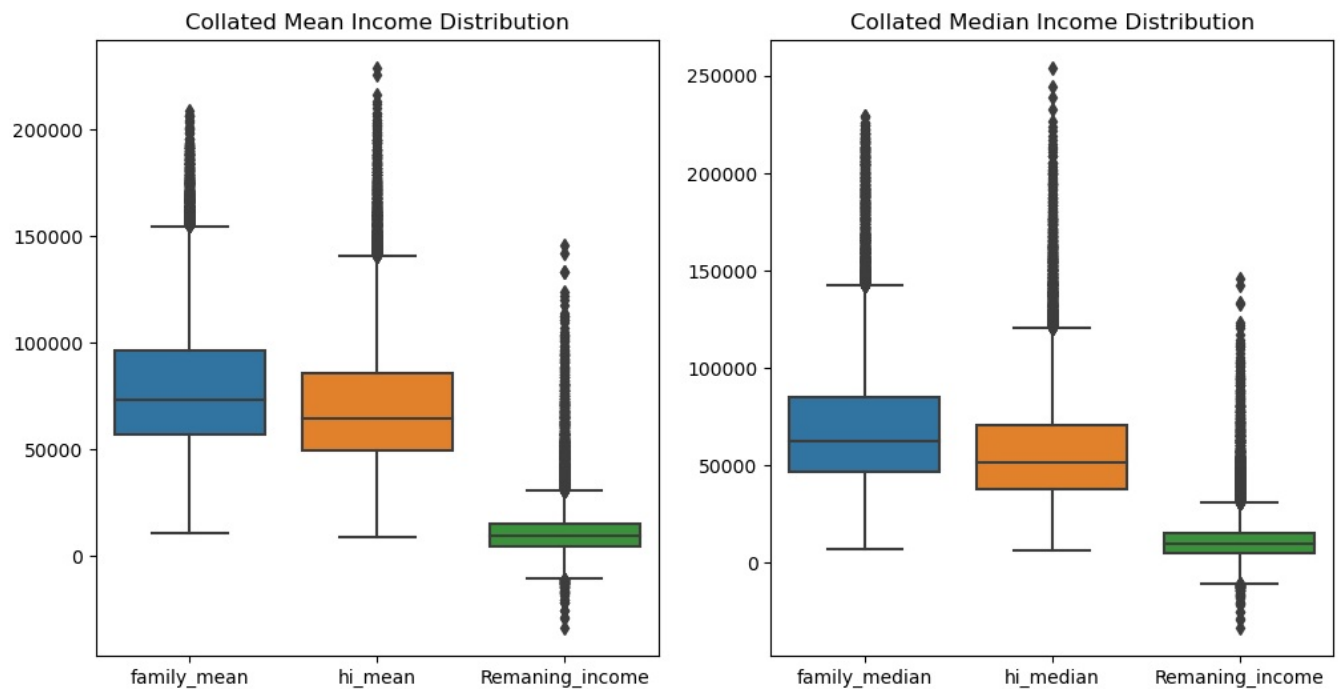
```
In [40]: df_train['Remaning_income'] = df_train['family_median'] - df_train['hi_median']
df_train
```

```
Out[40]:
```

	COUNTYID	STATEID	state	state_ab	city	place	type	zip_code	area_code	lat	...	pct_own	married
UID													
267822	53	36	New York	NY	Hamilton	Hamilton	City	13346	315	42.840812	...	0.79046	0.57851
246444	141	18	Indiana	IN	South Bend	Roseland	City	46616	574	41.701441	...	0.52483	0.34886
245683	63	18	Indiana	IN	Danville	Danville	City	46122	317	39.792202	...	0.85331	0.64745
279653	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	927	787	18.396103	...	0.65037	0.47257
247218	161	20	Kansas	KS	Manhattan	Manhattan City	City	66502	785	39.195573	...	0.13046	0.12356
...
279212	43	72	Puerto Rico	PR	Coamo	Coamo	Urban	769	787	18.076060	...	0.60422	0.24603
277856	91	42	Pennsylvania	PA	Blue Bell	Blue Bell	Borough	19422	215	40.158138	...	0.68072	0.61127
233000	87	8	Colorado	CO	Weldona	Saddle Ridge	City	80653	970	40.410316	...	0.78508	0.70451
287425	439	48	Texas	TX	Colleyville	Colleyville City	Town	76034	817	32.904866	...	0.93970	0.75503
265371	3	32	Nevada	NV	Las Vegas	Paradise	City	89123	702	36.064754	...	0.27912	0.34426


26585 rows × 81 columns

```
In [41]: plt.figure(figsize=(12,6))
plt.subplots_adjust(wspace=0.20)
plt.subplot(1,2,1)
sns.boxplot(data = df_train[['family_mean', 'hi_mean', 'Remaning_income']])
plt.title("Collated Mean Income Distribution")
plt.subplot(1,2,2)
sns.boxplot(data = df_train[['family_median', 'hi_median', 'Remaning_income']])
plt.title("Collated Median Income Distribution")
plt.show()
```



```
In [42]: plt.figure(figsize=(18,4))
plt.subplots_adjust(wspace=0.20)
plt.subplot(1,3,1)
sns.distplot(df_train['family_median'])
plt.title("Collated Family Income Distribution")
plt.subplot(1,3,2)
sns.distplot(df_train['hi_median'])
plt.title("Collated Household Income Distribution")
plt.subplot(1,3,3)
sns.distplot(df_train['Remaning_income'])
```

```
plt.title("Collated Remaining Income Distribution")
plt.show()
```




Week 2


Exploratory Data Analysis (EDA):

1. Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):

```
In [43]: fig, (ax1, ax2, ax3) = plt.subplots(3,1)
sns.distplot(df_train['pop'], ax = ax1)
sns.distplot(df_train['male_pop'], ax = ax2)
sns.distplot(df_train['female_pop'], ax = ax3)
plt.subplots_adjust(hspace=0.8)
plt.show()
```



```
In [44]: fig, (ax1, ax2) = plt.subplots(2,1)
sns.distplot(df_train['male_age_median'], ax = ax1)
sns.distplot(df_train['female_age_median'], ax = ax2)
plt.subplots_adjust(hspace=0.30)
plt.show()
```



a) Use pop and ALand variables to create a new field called population density

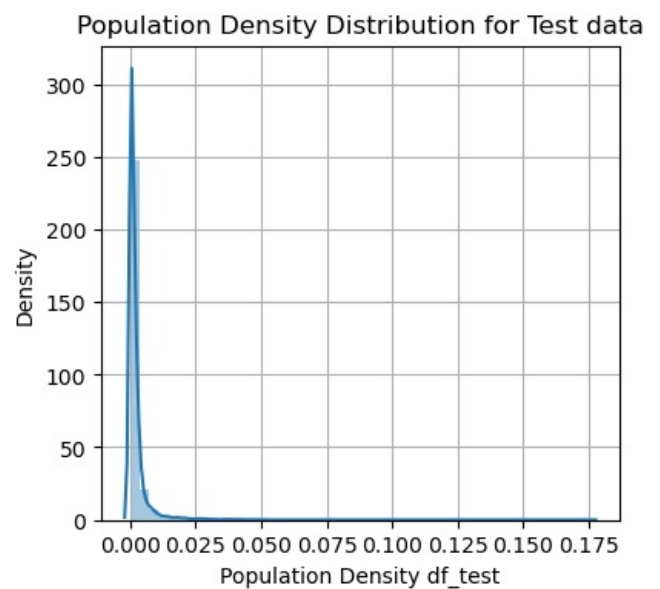
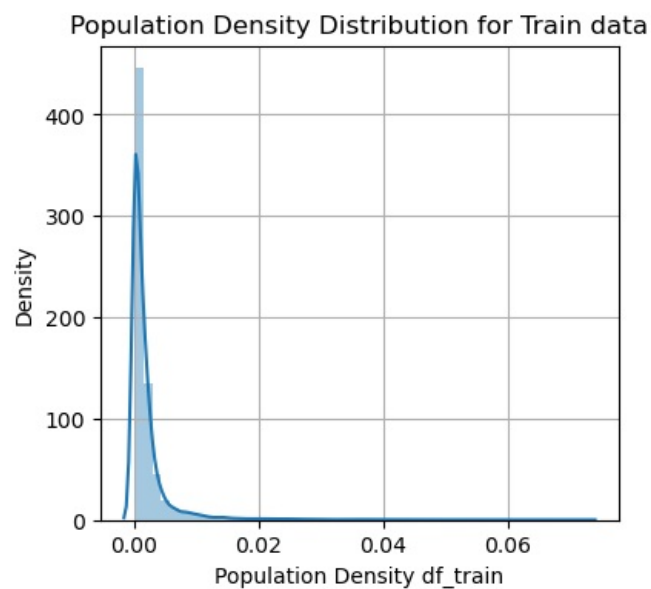
```
In [45]: df_train['pop_density'] = df_train['pop']/df_train['ALand']
df_test['pop_density'] = df_test['pop']/df_test['ALand']
df_train.head()
```

```
Out[45]:
```

	COUNTYID	STATEID	state	state_ab	city	place	type	zip_code	area_code	lat	...	married	married_snp	sep
	UID													
267822	53	36	New York	NY	Hamilton	Hamilton	City	13346	315	42.840812	...	0.57851	0.01882	0
246444	141	18	Indiana	IN	South Bend	Roseland	City	46616	574	41.701441	...	0.34886	0.01426	0
245683	63	18	Indiana	IN	Danville	Danville	City	46122	317	39.792202	...	0.64745	0.02830	0
279653	127	72	Puerto Rico	PR	San Juan	Guaynabo	Urban	927	787	18.396103	...	0.47257	0.02021	0
247218	161	20	Kansas	KS	Manhattan	Manhattan City	City	66502	785	39.195573	...	0.12356	0.00000	0

5 rows × 82 columns

```
In [46]: plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
sns.distplot(df_train['pop_density'])
plt.title("Population Density Distribution for Train data")
plt.xlabel("Population Density df_train")
plt.grid()
plt.subplot(1,2,2)
sns.distplot(df_test['pop_density'])
plt.title("Population Density Distribution for Test data")
plt.xlabel("Population Density df_test")
plt.grid()
plt.subplots_adjust(wspace= 0.30)
plt.show()
```



b) Use male_age_median, female_age_median, male_pop, and female_pop to create a new field called median age

```
In [47]: df_train['Median_age'] = (df_train['male_age_median'] * df_train['male_pop'] +
                                   df_train['female_age_median'] * df_train['female_pop'])/df_train['pop']
df_train[['pop', 'Median_age']].head()
```

```
Out[47]:
```

	pop	Median_age
UID		
267822	5230	44.667430
246444	2633	34.722748
245683	6881	41.774472
279653	2700	49.879012
247218	5637	21.965629


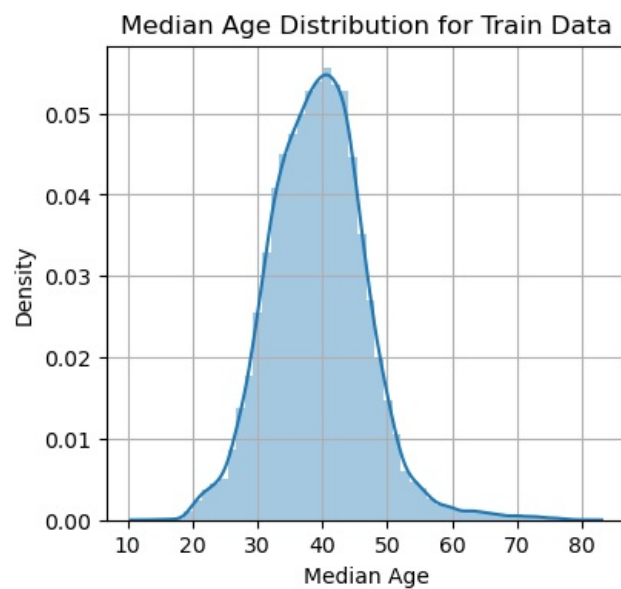
```
In [48]: df_test['Median_age'] = (df_test['male_age_median'] * df_test['male_pop'] +
                                   df_test['female_age_median'] * df_test['female_pop'])/df_test['pop']
df_test[['pop', 'Median_age']].head()
```

```
Out[48]:
```

	pop	Median_age
UID		
255504	3417	31.189053
252676	3796	46.382991
276314	3944	43.147420
248614	2508	45.155104
286865	6230	43.235983



c) Visualize the findings using appropriate chart type

```
In [49]: plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
sns.distplot(df_train['Median_age'])
plt.title("Median Age Distribution for Train Data")
plt.xlabel("Median Age")
plt.grid()
plt.subplot(1,2,2)
sns.distplot(df_test['Median_age'])
plt.title("Median Age Distribution for Test Data")
plt.xlabel("Median Age")
plt.grid()
plt.subplots_adjust(wspace= 0.30)
plt.show()
```



- Age ranges from 18 years to 75 years.
- The majority belongs to 40 years.
- Little right skewness is observed.

```
In [50]: plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
sns.boxplot(df_train['Median_age'])
plt.title("Population Density for Train Data")
plt.xlabel("Median Age")
plt.subplot(1,2,2)
sns.boxplot(df_test['Median_age'])
plt.title("Population Density for Test Data")
plt.xlabel("Median Age")
plt.subplots_adjust(wspace= 0.30)
plt.show()
```



2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.

```
In [51]: df_train['pop'].describe()
```

```
Out[51]: count    26585.000000
mean      4367.763438
std       2093.787018
min        63.000000
25%       2938.000000
50%       4078.000000
75%       5456.000000
max       53812.000000
Name: pop, dtype: float64
```

```
In [52]: df_train['pop_bins'] = pd.cut(df_train['pop'], bins = 5, labels = ['Very Low', 'Low', 'Medium', 'High', 'Very H:
```

```
In [53]: df_test['pop_bins'] = pd.cut(df_test['pop'], bins = 5, labels = ['Very Low', 'Low', 'Medium', 'High', 'Very Higl
```

```
In [54]: df_train[['pop', 'pop_bins']]
```

Out[54]:

	pop	pop_bins
UID		
267822	5230	Very Low
246444	2633	Very Low
245683	6881	Very Low
279653	2700	Very Low
247218	5637	Very Low
...
279212	1847	Very Low
277856	4155	Very Low
233000	2829	Very Low
287425	11542	Low
265371	3726	Very Low

26585 rows × 2 columns

```
In [55]: df_test[['pop', 'pop_bins']]
```

Out[55]:

	pop	pop_bins
UID		
255504	3417	Very Low
252676	3796	Very Low
276314	3944	Very Low
248614	2508	Very Low
286865	6230	Low
...
238088	5611	Very Low
242811	2695	Very Low
250127	7392	Low
241096	5945	Low
287763	4117	Very Low

11355 rows × 2 columns

```
In [56]: df_train.pop_bins.value_counts()
```

Out[56]:

Very Low	26334
Low	238
Medium	9
High	3
Very High	1

Name: pop_bins, dtype: int64

```
In [57]: df_test.pop_bins.value_counts()
```

Out[57]:

Very Low	8977
Low	2306
Medium	58
High	11
Very High	3

Name: pop_bins, dtype: int64

a) Analyze the married, separated, and divorced population for these population brackets

```
In [58]: df_train.groupby(by = 'pop_bins')[['married', 'separated', 'divorced']].count()
```

```
Out[58]:
```

	married	separated	divorced
pop_bins			
Very Low	26334	26334	26334
Low	238	238	238
Medium	9	9	9
High	3	3	3
Very High	1	1	1

```
In [59]: df_train.groupby(by = 'pop_bins')[['married', 'separated', 'divorced']].agg(['mean', 'median'])
```

```
Out[59]:
```

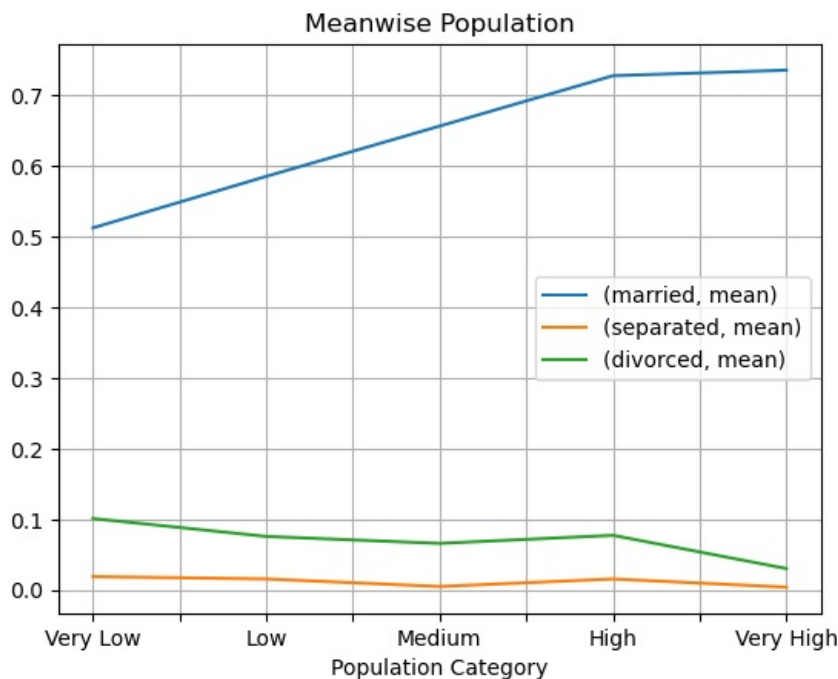
pop_bins	married		separated		divorced	
	mean	median	mean	median	mean	median
Very Low	0.511869	0.528105	0.019017	0.01351	0.101048	0.096090
Low	0.584691	0.592575	0.015655	0.01106	0.075749	0.070565
Medium	0.655737	0.618710	0.005003	0.00412	0.065927	0.064890
High	0.726957	0.736060	0.015663	0.00916	0.077310	0.063050
Very High	0.734740	0.734740	0.004050	0.00405	0.030360	0.030360

- In the Very Low population group, there are more divorced people as compared to the rest of the population category.
- Very High population group has highest number of married people, lowest number of separated as well as lowest number of divorced people as compared to the rest of the population category.

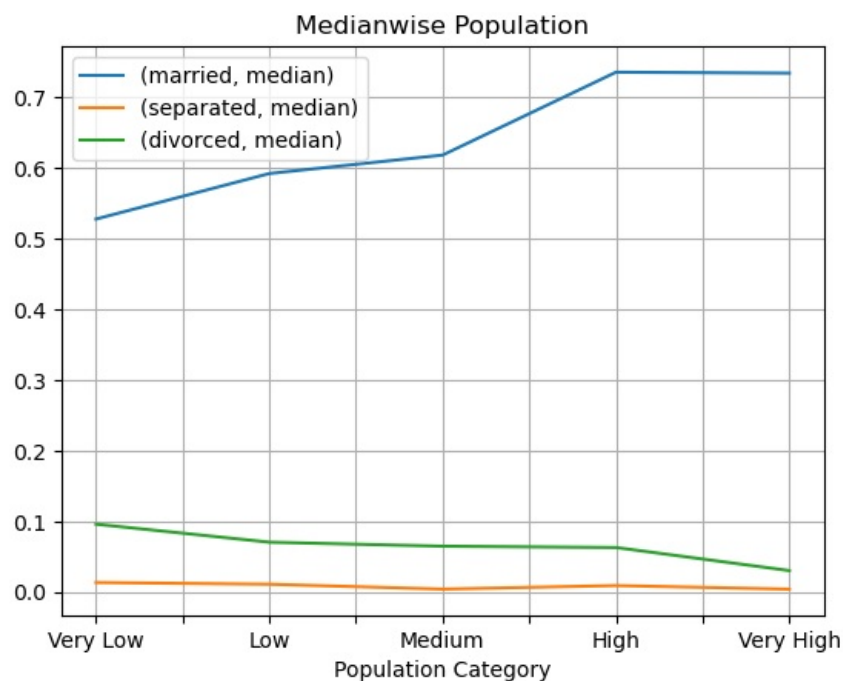
b) Visualize using appropriate chart type

```
In [60]: pop_bin_mean = df_train.groupby(by = 'pop_bins')[['married', 'separated', 'divorced']].agg(['mean'])
pop_bin_median = df_train.groupby(by = 'pop_bins')[['married', 'separated', 'divorced']].agg(['median'])
```

```
In [61]: pop_bin_mean.plot(title = "Meanwise Population")
plt.xlabel("Population Category")
plt.legend(loc = "best")
plt.grid()
plt.show()
```



```
In [62]: pop_bin_median.plot(title = "Medianwise Population")
plt.xlabel("Population Category")
plt.legend(loc = "best")
plt.grid()
plt.show()
```

3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.

```
In [63]: df_train[['state', 'rent_mean', 'rent_median']]
```

```
Out[63]:
```

	state	rent_mean	rent_median
UID			
267822	New York	769.38638	784.0
246444	Indiana	804.87924	848.0
245683	Indiana	742.77365	703.0
279653	Puerto Rico	803.42018	782.0
247218	Kansas	938.56493	881.0
...
279212	Puerto Rico	439.42839	419.0
277856	Pennsylvania	1813.19253	1788.0
233000	Colorado	849.39107	834.0
287425	Texas	1972.45746	1843.0
265371	Nevada	949.84199	924.0

26585 rows × 3 columns

```
In [64]: rent_state = df_train.groupby(by = 'state')['rent_mean'].mean()
rent_state.head()
```

```
Out[64]: state
Alabama      768.810406
Alaska       1173.830410
Arizona      1101.133798
Arkansas     715.367386
California   1479.363998
Name: rent_mean, dtype: float64
```

```
In [65]: income_state = df_train.groupby(by = 'state')['family_mean'].mean()
income_state.head()
```

```
Out[65]: state
Alabama      66814.665178
Alaska       92504.826703
Arizona      73546.551858
Arkansas     64046.416919
California   88438.468548
Name: family_mean, dtype: float64
```

```
In [66]: overall_percentage = rent_state/income_state
overall_percentage.head()
```

```
Out[66]: state
Alabama      0.011507
Alaska       0.012689
Arizona      0.014972
Arkansas     0.011170
California   0.016728
dtype: float64
```

```
In [67]: # Overall percentage
(df_train['rent_mean'].sum()/df_train['family_mean'].sum())*100
```

```
Out[67]: 1.329892795119824
```


4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

```
In [68]: df_train.columns
```

```
Out[68]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
               'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater', 'pop',
               'male_pop', 'female_pop', 'rent_mean', 'rent_median', 'rent_stdev',
               'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15',
               'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40',
               'rent_gt_50', 'universe_samples', 'used_samples', 'hi_mean',
               'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
               'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
               'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
               'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
               'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
               'home_equity_second mortgage', 'second mortgage', 'home equity', 'debt',
               'second mortgage_cdf', 'home equity_cdf', 'debt_cdf', 'hs_degree',
               'hs_degree_male', 'hs_degree_female', 'male_age_mean',
               'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
               'male_age_samples', 'female_age_mean', 'female_age_median',
               'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
               'pct_own', 'married', 'married_snp', 'separated', 'divorced',
               'BAD_DEBT', 'BINS_OD', 'BINS_BD', 'GOOD_DEBT', 'Remaning_income',
               'pop_density', 'Median_age', 'pop_bins'],
              dtype='object')
```

```
In [69]: cols_heatmap = df_train[['COUNTYID', 'STATEID', 'state', 'city', 'type', 'zip_code', 'pop', 'male_pop', 'female
                                'rent_mean', 'hi_mean', 'family_mean', 'second mortgage', 'home equity', 'debt', 'hs_deg
                                'male_age_mean', 'female_age_mean', 'pct_own', 'married', 'separated', 'divorced'],].co
```

```
In [70]: plt.figure(figsize=(25,15))
sns.heatmap(cols_heatmap, annot=True, cmap='YlGnBu')
plt.show()
```



- Very High Positive Correlation is observed between
 - male_pop, female_pop & pop
 - family_mean & hi_mean
 - family_age_mean & male_age_mean
- High Positive Correlation is observed between
 - hi_mean & rent_mean
 - family_mean & rent_mean
 - hs_degree & family_mean
 - married & pct_own

Project Task: Week 3

Data Pre-processing:

1. The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables.

2. Each variable is assumed to be dependent upon a linear combination of the common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as “specific variance” because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data. Following are the list of latent variables:

- Highschool graduation rates
- Median population age
- Second mortgage statistics
- Percent own
- Bad debt expense

```
In [71]: from factor_analyzer import FactorAnalyzer
```

```
In [72]: df_train.columns
```

```
Out[72]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
               'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater', 'pop',
               'male_pop', 'female_pop', 'rent_mean', 'rent_median', 'rent_stdev',
               'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15',
               'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40',
               'rent_gt_50', 'universe_samples', 'used_samples', 'hi_mean',
               'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
               'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
               'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
               'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
               'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
               'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
               'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
               'hs_degree_male', 'hs_degree_female', 'male_age_mean',
               'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
               'male_age_samples', 'female_age_mean', 'female_age_median',
               'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
               'pct_own', 'married', 'married_snp', 'separated', 'divorced',
               'BAD_DEBT', 'BINS_OD', 'BINS_BD', 'GOOD_DEBT', 'Remaning_income',
               'pop_density', 'Median_age', 'pop_bins'],
              dtype='object')
```

```
In [73]: fa_df_train = df_train.select_dtypes(exclude=('category', 'object'))
```

```
In [74]: fa_df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 26585 entries, 267822 to 265371
Data columns (total 76 columns):
```

#	Column	Non-Null Count	Dtype
0	COUNTYID	26585 non-null	int64
1	STATEID	26585 non-null	int64
2	zip_code	26585 non-null	int64
3	area_code	26585 non-null	int64
4	lat	26585 non-null	float64
5	lng	26585 non-null	float64
6	ALand	26585 non-null	float64
7	AWater	26585 non-null	int64
8	pop	26585 non-null	int64
9	male_pop	26585 non-null	int64
10	female_pop	26585 non-null	int64
11	rent_mean	26585 non-null	float64
12	rent_median	26585 non-null	float64
13	rent_stdev	26585 non-null	float64
14	rent_sample_weight	26585 non-null	float64
15	rent_samples	26585 non-null	float64
16	rent_gt_10	26585 non-null	float64
17	rent_gt_15	26585 non-null	float64
18	rent_gt_20	26585 non-null	float64
19	rent_gt_25	26585 non-null	float64
20	rent_gt_30	26585 non-null	float64
21	rent_gt_35	26585 non-null	float64
22	rent_gt_40	26585 non-null	float64
23	rent_gt_50	26585 non-null	float64
24	universe_samples	26585 non-null	int64
25	used_samples	26585 non-null	int64
26	hi_mean	26585 non-null	float64
27	hi_median	26585 non-null	float64
28	hi_stdev	26585 non-null	float64
29	hi_sample_weight	26585 non-null	float64
30	hi_samples	26585 non-null	float64
31	family_mean	26585 non-null	float64
32	family_median	26585 non-null	float64
33	family_stdev	26585 non-null	float64
34	family_sample_weight	26585 non-null	float64
35	family_samples	26585 non-null	float64
36	hc_mortgage_mean	26585 non-null	float64
37	hc_mortgage_median	26585 non-null	float64
38	hc_mortgage_stdev	26585 non-null	float64
39	hc_mortgage_sample_weight	26585 non-null	float64
40	hc_mortgage_samples	26585 non-null	float64
41	hc_mean	26585 non-null	float64
42	hc_median	26585 non-null	float64
43	hc_stdev	26585 non-null	float64
44	hc_samples	26585 non-null	float64
45	hc_sample_weight	26585 non-null	float64
46	home_equity_second_mortgage	26585 non-null	float64
47	second_mortgage	26585 non-null	float64
48	home_equity	26585 non-null	float64
49	debt	26585 non-null	float64
50	second_mortgage_cdf	26585 non-null	float64
51	home_equity_cdf	26585 non-null	float64
52	debt_cdf	26585 non-null	float64
53	hs_degree	26585 non-null	float64
54	hs_degree_male	26585 non-null	float64
55	hs_degree_female	26585 non-null	float64
56	male_age_mean	26585 non-null	float64
57	male_age_median	26585 non-null	float64
58	male_age_stdev	26585 non-null	float64
59	male_age_sample_weight	26585 non-null	float64
60	male_age_samples	26585 non-null	float64
61	female_age_mean	26585 non-null	float64
62	female_age_median	26585 non-null	float64
63	female_age_stdev	26585 non-null	float64
64	female_age_sample_weight	26585 non-null	float64
65	female_age_samples	26585 non-null	float64
66	pct_own	26585 non-null	float64
67	married	26585 non-null	float64
68	married_snp	26585 non-null	float64
69	separated	26585 non-null	float64
70	divorced	26585 non-null	float64
71	BAD_DEBT	26585 non-null	float64
72	GOOD_DEBT	26585 non-null	float64
73	Remaning_income	26585 non-null	float64
74	pop_density	26585 non-null	float64
75	Median_age	26585 non-null	float64

```
dtypes: float64(66), int64(10)
```

```
memory usage: 15.6 MB
```

```
In [75]: #Creating FactorAnalyzer object and performing factor analysis
```

```
fa = FactorAnalyzer(n_factors=25, rotation=None)
fa.fit_transform(fa_df_train)
loadings = fa.loadings_
```

```
#Checking Eigenvalues
```

```
ev, v = fa.get_eigenvalues()
ev
```

```
Out[75]: array([ 1.64501148e+01,  1.24030533e+01,  8.97275988e+00,  4.82727495e+00,
 4.41962370e+00,  3.38627145e+00,  2.37428417e+00,  2.13727698e+00,
 1.55751411e+00,  1.36381406e+00,  1.35566796e+00,  1.24811683e+00,
 1.14947852e+00,  1.05550496e+00,  1.00038367e+00,  9.07628682e-01,
 8.94817155e-01,  8.14675358e-01,  7.65526716e-01,  6.77394083e-01,
 6.05864840e-01,  5.87074096e-01,  5.80551737e-01,  5.56982551e-01,
 5.02233653e-01,  4.41424562e-01,  4.32444174e-01,  3.80558847e-01,
 3.47430797e-01,  3.13759836e-01,  3.07917477e-01,  2.96755983e-01,
 2.48148044e-01,  2.41685681e-01,  2.31477964e-01,  2.04300234e-01,
 1.97722934e-01,  1.76401829e-01,  1.60281026e-01,  1.51455296e-01,
 1.39304980e-01,  1.32498515e-01,  1.14040925e-01,  1.07609203e-01,
 1.03411124e-01,  9.13602308e-02,  9.06169709e-02,  7.99567607e-02,
 6.03726669e-02,  5.82364639e-02,  5.27144553e-02,  3.38775446e-02,
 3.17521110e-02,  2.81266062e-02,  2.27876143e-02,  2.10862680e-02,
 1.91096745e-02,  1.61396717e-02,  1.52777975e-02,  1.42064289e-02,
 1.17953170e-02,  8.22671465e-03,  6.61210078e-03,  5.38723042e-03,
 3.98041349e-03,  3.57107401e-03,  9.52446172e-04,  7.43395646e-04,
 5.92393085e-04,  3.43997631e-16,  2.66265617e-16,  2.26121051e-16,
 1.08834830e-16,  5.73952315e-17, -1.35508288e-16, -1.55325980e-16])
```

```
In [76]: print("Sorted:",sorted(ev, reverse=True))
print("\nSize:",ev.size)
print("\nSize:",fa_df_train.shape[1])
```


```
Sorted: [16.450114838545417, 12.40305327737094, 8.97275987782638, 4.827274951210276, 4.419623696403794, 3.3862714520271235, 2.374284165973279, 2.1372769776442353, 1.5575141141043185, 1.3638140600039537, 1.3556679565518985, 1.2481168257496817, 1.1494785225148565, 1.0555049624856776, 1.000383669709866, 0.9076286824534584, 0.8948171548744023, 0.8146753582956832, 0.7655267160358757, 0.677394083326619, 0.6058648396198004, 0.5870740961474066, 0.5805517365962005, 0.5569825511053841, 0.5022336533490541, 0.4414245620008767, 0.43244417359095744, 0.3805588472113495, 0.34743079715819264, 0.31375983577529487, 0.3079174765280145, 0.2967559828340237, 0.2481480441346208, 0.24168568055860046, 0.23147796370530083, 0.20430023400679853, 0.19772293437052163, 0.1764018288102684, 0.16028102638291536, 0.15145529649502545, 0.13930498047410328, 0.13249851450801228, 0.11404092486141937, 0.1076092025309672, 0.10341112380139938, 0.09136023075056879, 0.09061697093655742, 0.0799567606555546, 0.060372666910959064, 0.05823646386647841, 0.05271445525119728, 0.033877544638451634, 0.0317521110232445, 0.028126606150349686, 0.02278761429873255, 0.02108626799074645, 0.019109674520868977, 0.01613967168761339, 0.015277797503950728, 0.014206428922813057, 0.011795316972700639, 0.008226714650811862, 0.006612100779063703, 0.005387230420946085, 0.00398041348923252, 0.003571074012018963, 0.0009524461722090671, 0.0007433956461345817, 0.0005923930845492011, 3.439976305463639e-16, 2.662656174824726e-16, 2.2612105112295604e-16, 1.0883483011441228e-16, 5.739523147963836e-17, -1.355082876556648e-16, -1.553259804951645e-16]
```

```
Size: 76
```

```
Size: 76
```

```
In [77]: xval = range(1, fa_df_train.shape[1]+1)
```

```
In [78]: plt.figure(figsize=(15,10))
plt.title("Scree plot")
plt.xlabel("Factors")
plt.ylabel("Eigen Values")
plt.scatter(xval, ev)
plt.plot(xval, ev)
plt.grid(b=True, which='minor', color='r', linestyle='--')
plt.minorticks_on()
plt.show()
```



```
In [79]: df_Factors = pd.DataFrame.from_records(loadings)

df_Factors = df_Factors.add_prefix('Factor ')

df_Factors.index = fa_df_train.columns
df_Factors
```

```
Out[79]:
```

	Factor 0	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	Factor 8	Factor 9	...	Factor 15
COUNTYID	-0.121467	0.023864	0.064159	-0.049338	0.043508	-0.072132	-0.064150	-0.000500	0.019670	0.083287	...	-0.041960
STATEID	-0.116626	-0.025949	0.174363	-0.084042	0.098171	0.007296	-0.386308	-0.106417	0.052836	0.008942	...	0.140697
zip_code	-0.041975	0.108943	-0.068539	-0.144379	-0.033195	-0.202175	0.749217	0.441724	0.114256	-0.109741	...	0.181417
area_code	0.014956	0.020605	-0.011753	0.024296	-0.005710	-0.084730	0.015043	-0.055767	-0.037316	0.035503	...	-0.009654
lat	0.189045	-0.105997	-0.067375	-0.175523	-0.049946	0.186421	-0.188249	0.083182	0.126231	-0.245474	...	0.238567
...
BAD_DEBT	0.650683	0.029018	-0.323422	-0.111934	-0.461837	0.303323	-0.002725	-0.166413	0.066974	-0.127652	...	0.040025
GOOD_DEBT	-0.650683	-0.029018	0.323422	0.111934	0.461837	-0.303323	0.002725	0.166413	-0.066974	0.127652	...	-0.040025
Remaning_income	0.353804	-0.157782	0.001032	0.159686	0.306825	0.286105	-0.112850	0.184142	0.283166	0.064099	...	-0.210747
pop_density	-0.009890	0.084188	-0.413661	0.098072	0.328215	0.082651	0.008607	-0.192113	-0.067117	0.021249	...	-0.042315
Median_age	0.320412	-0.377575	0.616295	0.401649	-0.046410	0.313298	0.185335	-0.067104	-0.180255	0.069782	...	0.004953

76 rows × 25 columns

```
In [80]: fa = FactorAnalyzer(n_factors=12 ,rotation="varimax")
fa.fit(fa_df_train)
new_load = fa.loadings_
```

```
In [81]: New_Factors = pd.DataFrame.from_records(new_load)
New_Factors = New_Factors.add_prefix("Factor ")
New_Factors.index = fa_df_train.columns
New_Factors
```


Out[81]:		Factor 0	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	Factor 8	Factor 9	Factor 10	Factor 11
	COUNTYID	-0.094306	0.016231	-0.067259	-0.026929	-0.032105	-0.004703	-0.143706	-0.002368	0.048996	-0.008294	-0.049634	-0.000000
	STATEID	-0.115687	-0.002394	-0.065692	-0.032880	-0.094090	0.082233	-0.147470	-0.057465	0.335893	-0.047756	-0.088017	0.000000
	zip_code	-0.044881	0.049550	-0.105693	0.014041	-0.034214	0.003868	-0.069760	0.050855	-0.894686	-0.026567	0.012577	0.010000
	area_code	0.041746	0.030346	-0.017843	-0.054318	-0.004943	-0.075285	-0.030787	0.009855	0.004289	0.050287	-0.001366	-0.010000
	lat	0.079279	-0.110047	-0.077106	0.051019	-0.074011	0.251972	0.181171	0.046492	0.093402	-0.086824	0.079223	0.100000

	BAD_DEBT	0.361262	0.045864	-0.017889	-0.041039	-0.026821	0.153037	0.846767	0.311941	0.026378	0.066470	0.070511	-0.020000
	GOOD_DEBT	-0.361262	-0.045864	0.017889	0.041039	0.026821	-0.153037	-0.846767	-0.311941	-0.026378	-0.066470	-0.070511	0.020000
	Remaning_income	0.371282	-0.101728	0.158945	0.193463	0.025081	0.355310	0.014490	-0.055318	0.059092	-0.085131	-0.216458	0.000000
	pop_density	0.255545	-0.060181	-0.114827	0.393536	0.084184	-0.271423	-0.038532	-0.021890	0.105213	0.027958	0.061886	-0.110000
	Median_age	0.133890	-0.072269	0.946916	-0.202673	-0.070110	0.111837	0.026530	-0.036789	0.032941	-0.038426	-0.024761	0.110000

76 rows × 12 columns

In [82]: New_Factors.columns

Out[82]: Index(['Factor 0', 'Factor 1', 'Factor 2', 'Factor 3', 'Factor 4', 'Factor 5',
'Factor 6', 'Factor 7', 'Factor 8', 'Factor 9', 'Factor 10',
'Factor 11'],
dtype='object')

In [83]: New_Factors_df = round(New_Factors.loc[['hs_degree', 'hs_degree_male', 'hs_degree_female', "male_age_median", "female_age_median",
"home_equity_second_mortgage", 'second_mortgage', 'second_mortgage_cdf', 'pct_own

```
In [84]: def color_negative_red(value):
    """
    Colors elements in a dataframe
    green if positive and red if
    negative. Does not color NaN
    values.
    """

    if value < -0.6:
        color = 'red'
    elif value > 0.6:
        color = 'green'
    else:
        color = 'black'

    return 'color: %s' % color
```

In [85]: New_Factors_df.style.applymap(color_negative_red)

Out[85]:		Factor 0	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	Factor 8	Factor 9	Factor 10	Factor 11
	hs_degree	0.000000	0.000000	0.000000	-0.000000	-0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
	hs_degree_male	0.000000	0.000000	0.000000	-0.000000	-0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
	hs_degree_female	0.000000	0.000000	0.000000	-0.000000	-0.000000	1.000000	0.000000	0.000000	0.000000	-0.000000	0.000000	0.000000
	male_age_median	0.000000	-0.000000	1.000000	-0.000000	-0.000000	0.000000	0.000000	-0.000000	0.000000	-0.000000	-0.000000	-0.000000
	female_age_median	0.000000	-0.000000	1.000000	-0.000000	-0.000000	0.000000	0.000000	-0.000000	0.000000	-0.000000	-0.000000	-0.000000
	home_equity_second_mortgage	0.000000	0.000000	-0.000000	0.000000	0.000000	-0.000000	0.000000	1.000000	-0.000000	0.000000	0.000000	0.000000
	second_mortgage	0.000000	0.000000	-0.000000	0.000000	0.000000	-0.000000	0.000000	1.000000	-0.000000	0.000000	0.000000	0.000000
	second_mortgage_cdf	-0.000000	-0.000000	0.000000	0.000000	0.000000	-0.000000	-0.000000	-1.000000	0.000000	-0.000000	-0.000000	-0.000000
	pct_own	0.000000	0.000000	0.000000	-1.000000	-0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.000000	-0.000000
	BAD_DEBT	0.000000	0.000000	-0.000000	-0.000000	-0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000

We can see that 'Related parameters' are loading on Unique Factors

In [86]: len(fa_df_train.columns)

Out[86]: 76

```
In [87]: #Fetching the variance
fact_var = fa.get_factor_variance()
fact_var
```



```
Out[87]: (array([12.32196237, 11.91316203, 5.58056077, 5.49365386, 4.74170788,
3.78490369, 3.73084787, 2.75661873, 2.16353751, 1.90916053,
1.76345868, 1.54623546]),
array([0.16213108, 0.15675213, 0.07342843, 0.07228492, 0.06239089,
0.04980136, 0.0490901 , 0.0362713 , 0.0284676 , 0.02512053,
0.0232034 , 0.0203452 ]),
array([0.16213108, 0.31888322, 0.39231165, 0.46459657, 0.52698746,
0.57678882, 0.62587893, 0.66215023, 0.69061783, 0.71573836,
0.73894176, 0.75928697]))
```

```
In [88]: Factor_variance = pd.DataFrame.from_records(fact_var)
Factor_variance = Factor_variance.add_prefix("Factor ")
Factor_variance
```

```
Out[88]:
```

	Factor 0	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	Factor 8	Factor 9	Factor 10	Factor 11
0	12.321962	11.913162	5.580561	5.493654	4.741708	3.784904	3.730848	2.756619	2.163538	1.909161	1.763459	1.546235
1	0.162131	0.156752	0.073428	0.072285	0.062391	0.049801	0.049090	0.036271	0.028468	0.025121	0.023203	0.020345
2	0.162131	0.318883	0.392312	0.464597	0.526987	0.576789	0.625879	0.662150	0.690618	0.715738	0.738942	0.759287

```
In [89]: Factor_variance.index = ['Loadings', 'Proportion Var', 'Cummulative Var']
round(Factor_variance,2)
```

```
Out[89]:
```


	Factor 0	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	Factor 8	Factor 9	Factor 10	Factor 11
Loadings	12.32	11.91	5.58	5.49	4.74	3.78	3.73	2.76	2.16	1.91	1.76	1.55
Proportion Var	0.16	0.16	0.07	0.07	0.06	0.05	0.05	0.04	0.03	0.03	0.02	0.02
Cummulative Var	0.16	0.32	0.39	0.46	0.53	0.58	0.63	0.66	0.69	0.72	0.74	0.76

Project Task: Week 4

Data Modeling :

1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deplotment_RE.xlsx'. Column hc_mortgage_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location. Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc_mortgage_mean.
 - a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.
 - b) Run another model at State level. There are 52 states in USA.
 - c) Keep below considerations while building a linear regression model. Data Modeling :
 - Variables should have significant impact on predicting Monthly mortgage and owner costs
 - Utilize all predictor variable to start with initial hypothesis
 - R square of 60 percent and above should be achieved
 - Ensure Multi-collinearity does not exist in dependent variables
 - Test if predicted variable is normally distributed

```
In [90]: plt.figure(figsize=(10,7), dpi= 80)
sns.distplot(df_train.hc_mortgage_mean, color="dodgerblue", label="hc_mortgage_mean")
plt.legend();
plt.show()
```



Target variable has a positive skewness

```
In [91]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error, accuracy_score
```

```
In [92]: print(df_train.shape)
print(df_test.shape)
```

```
(26585, 84)
(11355, 79)
```

```
In [93]: test_list = df_test.columns
lis = []
for col in df_train:
    if col not in test_list:
        lis.append(col)
lis
```

```
Out[93]: ['BAD_DEBT', 'BINS_OD', 'BINS_BD', 'GOOD_DEBT', 'Remaning_income']
```

```
In [94]: print(df_train.shape)
df_train.drop(columns=['BAD_DEBT', 'BINS_OD', 'BINS_BD', 'GOOD_DEBT', 'Remaning_income'], inplace = True)
print(df_train.shape)
```

```
(26585, 84)
(26585, 79)
```

```
In [95]: print(df_train.shape)
print(df_test.shape)
```

```
(26585, 79)
(11355, 79)
```

```
In [96]: df_train[['hc_mortgage_mean']]
```

Out[96]:

hc_mortgage_mean	
UID	
267822	1414.80295
246444	864.41390
245683	1506.06758
279653	1175.28642
247218	1192.58759
...	...
279212	770.11560
277856	2210.84055
233000	1671.07908
287425	3074.83088
265371	1455.42340

26585 rows × 1 columns

In [97]:

```
df_train.dropna(axis = 0, inplace=True)
df_test.dropna(axis = 0, inplace=True)
```

In [98]:

```
print(df_train.shape)
print(df_test.shape)
```

(26585, 79)
(11355, 79)

In [99]:

```
feature_cols=['COUNTYID','STATEID','zip_code','pop', 'family_mean', 'second_mortgage', 'home_equity', 'debt','h
'Median_age','pct_own', 'married','separated', 'divorced']
```

In [100...]

```
week4_bck_train = df_train.copy()
week4_bck_test = df_test.copy()
```

In [101...]

```
x_train = df_train[feature_cols]
y_train = df_train['hc_mortgage_mean']
x_test = df_test[feature_cols]
y_test = df_test['hc_mortgage_mean']
```

In [102...]

```
from sklearn.preprocessing import StandardScaler
```

In [103...]

```
x_train.describe()
```

Out[103]:

	COUNTYID	STATEID	zip_code	pop	family_mean	second_mortgage	home_equity	debt	hs_degree
count	26585.000000	26585.000000	26585.000000	26585.000000	26585.00000	26585.000000	26585.000000	26585.000000	26585.000000
mean	85.580440	28.256348	50134.895091	4367.763438	79282.24354	0.029876	0.100928	0.629911	0.859577
std	97.891735	16.370924	29492.900596	2093.787018	31125.02069	0.030664	0.065590	0.148977	0.110139
min	1.000000	1.000000	602.000000	63.000000	10706.26180	0.000000	0.000000	0.013590	0.186520
25%	29.000000	13.000000	27106.000000	2938.000000	57217.14016	0.008060	0.050450	0.539480	0.809700
50%	63.000000	27.000000	47905.000000	4078.000000	73119.26510	0.022730	0.094900	0.648080	0.889340
75%	109.000000	42.000000	77084.000000	5456.000000	96218.70699	0.042900	0.143680	0.736190	0.939040
max	840.000000	72.000000	99925.000000	53812.000000	208969.99390	0.608700	0.687500	0.978260	1.000000

In [104...]

```
y_train.describe()
```

Out[104]:

count	26585.000000
mean	1627.898787
std	620.559056
min	402.681840
25%	1158.136460
50%	1459.286080
75%	1979.249110
max	4462.342290
Name: hc_mortgage_mean, dtype: float64	

In [105...]

```
sc = StandardScaler()
x_train_scaled = sc.fit_transform(x_train)
x_test_scaled =sc.fit_transform(x_test)
```

In [106...]

```
x_train_scaled
```

```
Out[106]: array([[ -0.33282741,  0.47302145, -1.24740484, ...,  0.5043254 ,
                -0.32824767, -0.27477163],
                [ 0.56614179, -0.62650957, -0.11931554, ..., -1.25295184,
                -0.23547174, -0.22025652],
                [-0.23067182, -0.62650957, -0.13606565, ...,  1.03185291,
                -0.14518977,  0.12088226],
                ...,
                [ 0.0145016 , -1.23736014,  1.03478052, ...,  1.46847489,
                -0.68738034, -0.49660619],
                [ 3.6103784 ,  1.20604213,  0.87816361, ...,  1.85505293,
                -0.49035617, -1.01051593],
                [-0.84360537,  0.22868122,  1.32197368, ..., -1.28815095,
                0.55212616,  0.67924277]])
```

```
In [107]: x_test_scaled
```

```
Out[107]: array([[ 0.77154843, -0.14787648, -0.06699701, ..., -1.71167277,
                0.94978465,  0.90847794],
                [-0.85528175, -0.32862972, -1.55088833, ...,  0.99225534,
                -0.95828531,  0.71431416],
                [-0.71469149,  0.81614082, -1.19158503, ...,  0.67232607,
                -0.27872617,  0.0086278 ],
                ...,
                [-0.77494446, -0.20812756, -1.63072977, ..., -0.05555059,
                -0.90424084, -0.56670348],
                [-0.59418555, -0.56963404,  0.03957057, ...,  1.17835505,
                -0.95828531, -1.11423691],
                [ 3.6837753 ,  1.17764731,  0.96113419, ...,  0.06859098,
                0.30275228,  0.12655808]])
```

Running a model at Nation Level.

```
In [108]: lr = LinearRegression()
```

```
In [109]: lr.fit(x_train_scaled, y_train)
```

```
Out[109]: LinearRegression()
```

```
In [110]: y_pred = lr.predict(x_test_scaled)
```

```
In [111]: print("Overall R2 Score", r2_score(y_test, y_pred))
          print("Overall RMSE", np.sqrt(mean_squared_error(y_test, y_pred)))
```

Overall R2 Score 0.7513061476377627

Overall RMSE 314.79945199484854

R2 Score is 75% and RMSE is 314 which is good but we will still proceed with running the model at State Level

Running a model at state level

```
In [112]: uni_state = df_train.STATEID.unique()
          uni_state[:10:2]
```

```
Out[112]: array([36, 72,  1, 45,  5], dtype=int64)
```

```
In [113]: for i in [20, 1, 45]:
          print("*****90)
          print("STATEID :", i)

          x_train_state = df_train[df_train.STATEID == i][feature_cols]
          y_train_state = df_train[df_train.STATEID == i]['hc_mortgage_mean']

          x_test_state = df_test[df_test.STATEID == i][feature_cols]
          y_test_state = df_test[df_test.STATEID == i]['hc_mortgage_mean']

          x_train_state_scaled = sc.fit_transform(x_train_state)
          x_test_state_scaled = sc.fit_transform(x_test_state)

          lr = LinearRegression()

          lr.fit(x_train_state_scaled, y_train_state)
          y_pred_state = lr.predict(x_test_state_scaled)

          print("Overall R2 Score for", i, ":", r2_score(y_test_state, y_pred_state))
          print("Overall RMSE for", i, ":", np.sqrt(mean_squared_error(y_test_state, y_pred_state)))
```

```

*****
STATEID : 20
Overall R2 Score for 20 : 0.8801114508360377
Overall RMSE for 20 : 143.41926392999306
*****

STATEID : 1
Overall R2 Score for 1 : 0.6729192403324948
Overall RMSE for 1 : 174.81015560089193
*****

STATEID : 45
Overall R2 Score for 45 : 0.7068418321829575
Overall RMSE for 45 : 181.52977719088605

```

Checking the Residuals

```

In [114]: residuals = y_test - y_pred
residuals

```

```


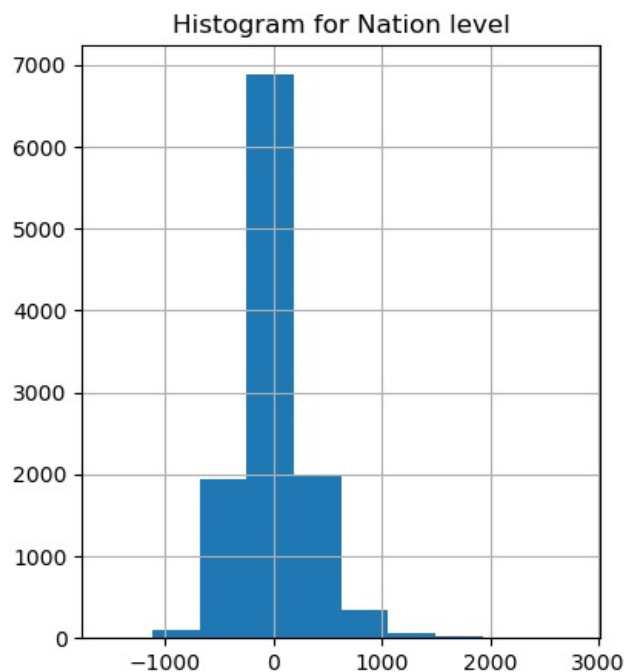
Out[114]: UID
255504    251.488994
252676    -7.849712
276314    107.468817
248614   -107.329935
286865   -54.941582
...
238088   -19.295706
242811  -163.964752
250127   -69.465187
241096  -274.083718
287763    22.932771
Name: hc_mortgage_mean, Length: 11355, dtype: float64

```

```

In [115]: plt.figure(figsize=(10,5))
plt.subplots_adjust(wspace=0.30)
plt.subplot(1,2,1)
plt.hist(residuals)
plt.grid()
plt.title("Histogram for Nation level")
plt.subplot(1,2,2)
plt.hist(y_test_state - y_pred_state)
plt.title("Histogram for State level")
plt.grid()
plt.show()

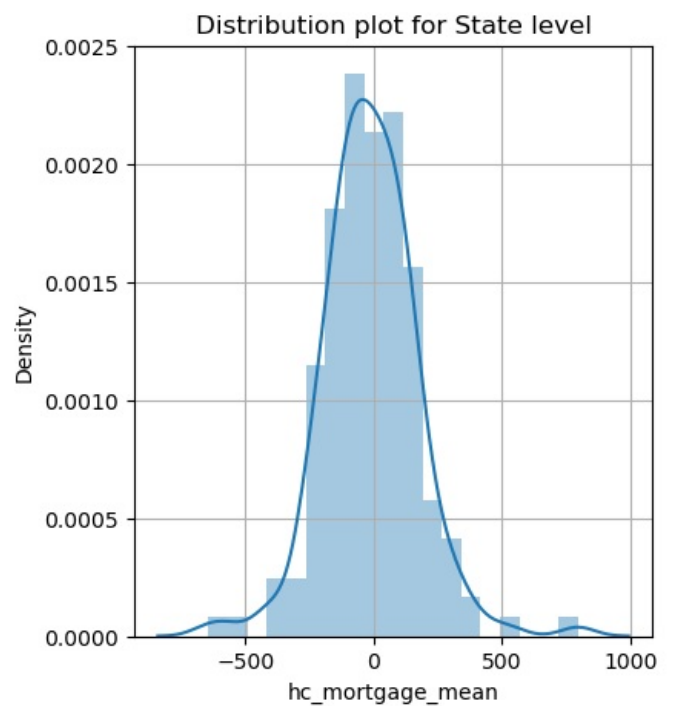
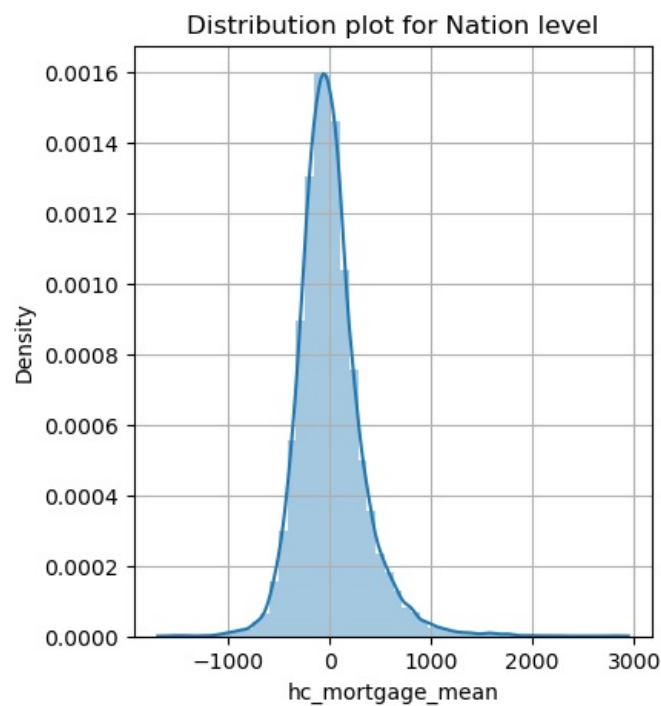
```



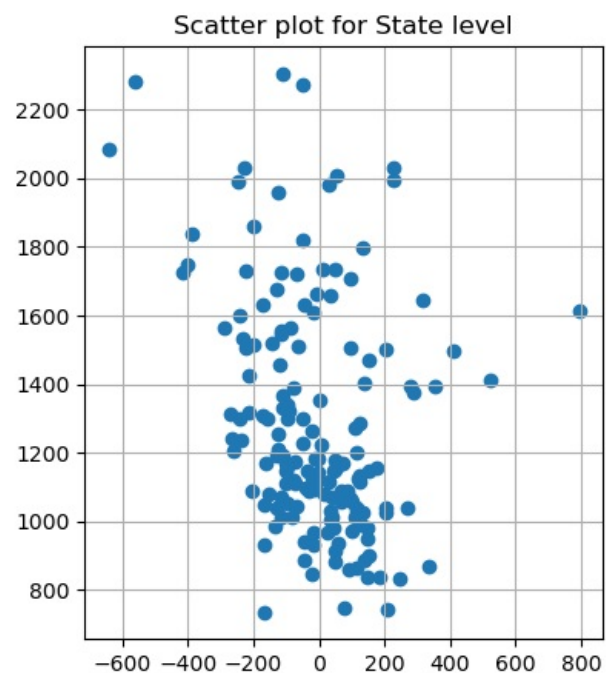
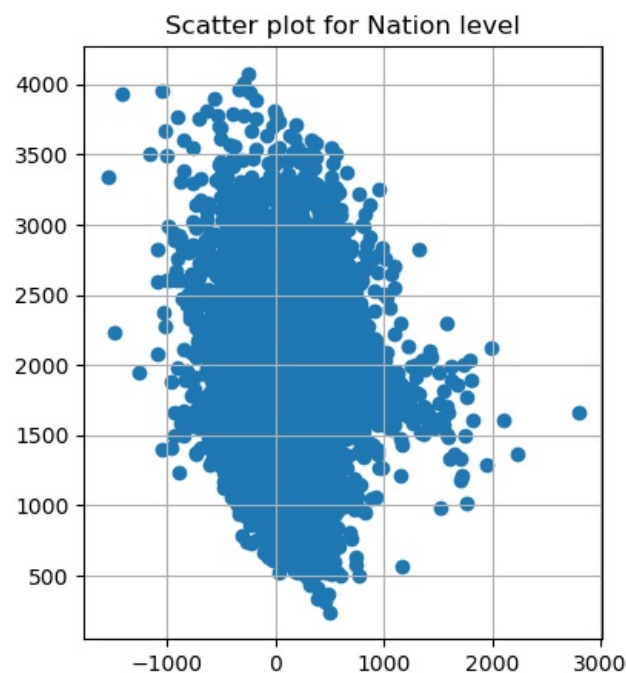
```

In [116]: plt.figure(figsize=(10,5))
plt.subplots_adjust(wspace=0.30)
plt.subplot(1,2,1)
sns.distplot(residuals)
plt.grid()
plt.title("Distribution plot for Nation level")
plt.subplot(1,2,2)
sns.distplot(y_test_state - y_pred_state)
plt.title("Distribution plot for State level")
plt.grid()
plt.show()

```





```
In [117]: plt.figure(figsize=(10,5))
plt.subplots_adjust(wspace=0.30)
plt.subplot(1,2,1)
plt.scatter(residuals, y_pred)
plt.grid()
plt.title("Scatter plot for Nation level")
plt.subplot(1,2,2)
plt.scatter((y_test_state - y_pred_state), y_pred_state)
plt.title("Scatter plot for State level")
plt.grid()
plt.show()
```



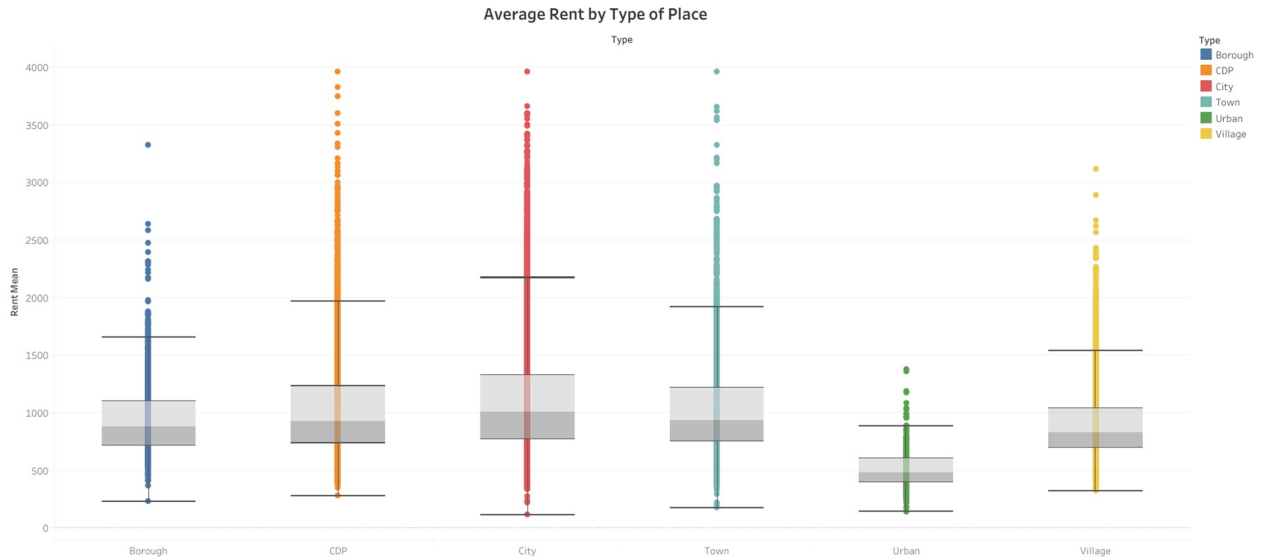
Data Reporting:

2. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:




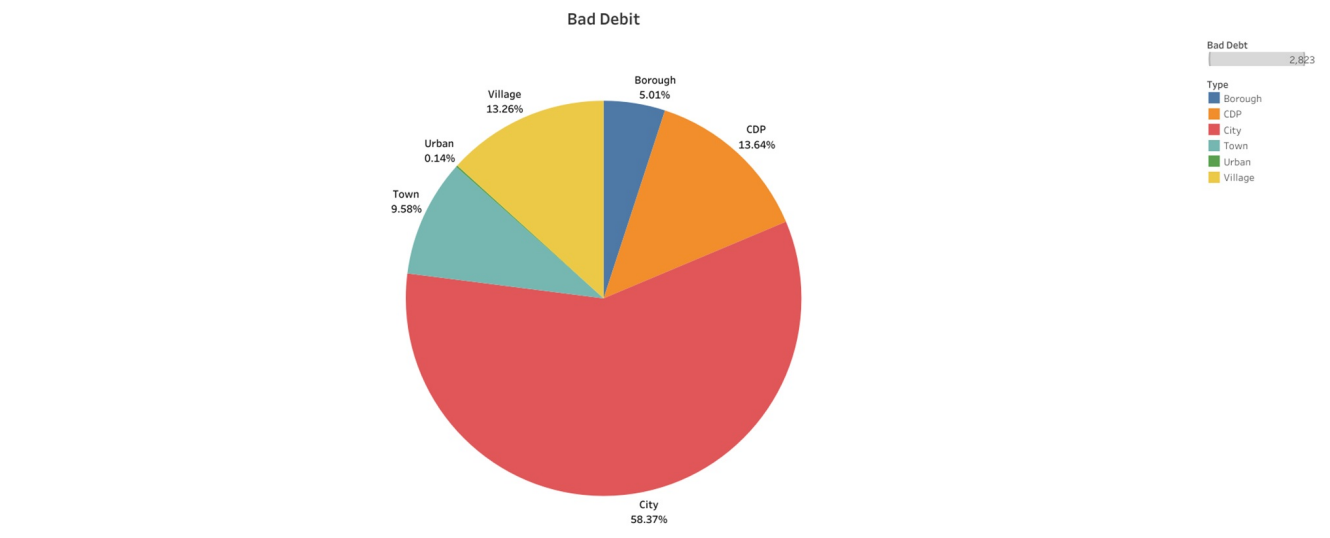


a) Box plot of distribution of average rent by type of place (village, urban, town, etc.).

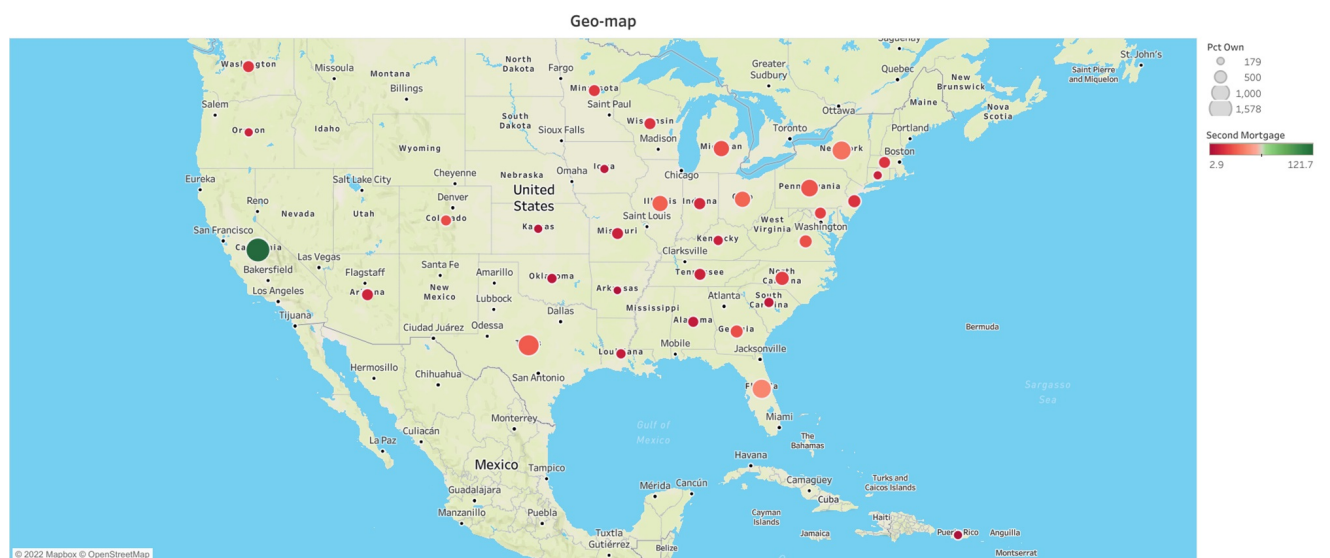


b) Pie charts to show overall debt and bad debt.







c) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map.




d) Heat map for correlation matrix.



e) Pie chart to show the population distribution across different types of places (village, urban, town etc.)





Loading [MathJax]/extensions/Safe.js